

Chapter 13: Union of Artificial Intelligence and Molecular Docking

Elvis Martis

1 Introduction

Classical docking algorithms combine a conformational search over ligand (and sometimes receptor) degrees of freedom with an empirical or physics-based scoring function intended to approximate the binding free energy landscape. Although these approaches have enabled decades of successful drug discovery programs, their accuracy is often limited by (over)simplified energy models, incomplete treatment of protein flexibility, and heuristic search schemes that struggle in rugged, high-dimensional free energy landscapes.

Over the past decade, artificial intelligence (AI) and deep learning methods have begun to transform the field of drug discovery, particularly molecular docking. Convolutional neural network (CNN)-based scoring functions, such as those in GNINA[1], have significantly improved pose ranking relative to AutoDock Vina, particularly when rescoring pre-generated pose ensembles. These methods learn complex, non-linear relationships between 3D protein-ligand complex and binding quality directly from structural data. This somewhat bypasses the need to fine-tune the functional forms of the scoring function manually[2]. More recently, generative and geometric deep learning models like *EquiBind*[3] and *DiffDock*[4] have started to replace or augment the classical conformational search itself, learning to propose plausible binding poses directly, often with large speedups over sampling-based methods.

Moreover, breakthroughs in protein structure prediction, most notably by AlphaFold2[5] and 3[6], have made high-quality structural models of previously intractable targets widely available. Building upon these ideas, new all-atom foundation models such as Boltz-2[7] can co-fold proteins, nucleic acids and ligands into complexes, while simultaneously predicting binding

affinity with accuracy approaching free-energy perturbation (FEP) at orders-of-magnitude lower computational cost.

2 Classical Molecular Docking: Search and Scoring

Given a protein receptor R and a ligand L , molecular docking seeks the bound complex configuration X^* that approximately minimises the binding free energy

$$\Delta G(R, L) = G(RL) - G(R) - G(L), \quad (1)$$

where $G(\cdot)$ denotes the free energy of the system. In practice, docking approximates this task by:

1. Sampling a discrete set of candidate poses $\{X_k\}_{k=1}^K$ over ligand translation, rotation and torsions. If the receptor is treated as flexible, then side-chain or backbone degrees of freedom are allowed to move.
2. Scoring each pose with an approximate energy or scoring function $S(X_k)$, selecting poses that score very high (as a convention, a more negative number is better) as predictions.

Classical scoring functions can be written generically as

$$S(X) = \sum_i w_i f_i(X) \quad (2)$$

where $f_i(X)$ are features, e.g., van der Waals interactions, hydrogen bonding, desolvation, electrostatics, rotatable bond penalties, and w_i are parameters obtained by fitting to experimental data. Search algorithms typically combine stochastic moves, such as Monte Carlo, genetic algorithms, Lamarckian GA, systematic search or any combination thereof, with local minimisation in the space of ligand pose and conformation.

Despite many refinements, three key challenges remain:

- Approximate scoring: Linear or low-order functional forms struggle to capture subtle enthalpy–entropy balance, water-mediated interactions, and long-range effects.

- Limited flexibility: Most workflows treat the receptor as rigid or allow only limited side-chain motion, missing induced-fit or allosteric rearrangements.
- Search complexity: Heuristic sampling can undersample near-native basins in large, flexible ligands or in cryptic pockets, leading to false negatives.

3 Basics of Deep Learning

3.1 Feedforward networks and loss minimisation

Deep learning methods learn a parametric function $f_\theta : \mathbf{x} \rightarrow \mathbf{y}$ from input features $\mathbf{x}$, e.g., a 3D protein–ligand representation, to outputs $\mathbf{y}$, e.g., binding affinity, pose quality, or coordinates, by minimising a loss function over a dataset of examples $\{(x^{(n)}, y^{(n)})\}_{n=1}^N$. A simple fully connected dense layer computes

$$h^{(\ell+1)} = \sigma(W^{(\ell)}h^{(\ell)} + b^{(\ell)}) \quad (3)$$

where $h^{(\ell)}$ is the input vector to layer ℓ , $W^{(\ell)}$ and $b^{(\ell)}$ are learnable weights and biases respectively, and σ is a non-linear activation function, such as ReLU[8], Sigmoid[9] or Tanh. For regression tasks like binding affinity prediction, the network outputs a scalar $\hat{y}$ and the standard loss is mean-squared error (MSE):

$$\mathcal{L}_{\text{MSE}}(\theta) = \frac{1}{N} \sum_{n=1}^N (\hat{y}^{(n)} - y^{(n)})^2 \quad (4)$$

For classification, such as binder vs non-binder, one uses a cross-entropy loss. Training proceeds by stochastic gradient descent (SGD) or variants (Adam, RMSProp)[10], using backpropagation to compute gradients $\partial \mathcal{L} / \partial \theta$ and update θ .

3.2 Convolutional and 3D voxel models

To model spatially local patterns in images or 3D densities, convolutional neural networks (CNNs) replace dense layers with convolutions. In 3D, a convolutional layer applies a learnable kernel K across a voxel grid representation of a protein–ligand complex:

$$(h * K)(\mathbf{r}) = \sum_{\mathbf{s}} h(\mathbf{r} - \mathbf{s})K(\mathbf{s}) \quad (5)$$

where $\mathbf{r}$ and $\mathbf{s}$ index grid positions in 3D space. CNNs are translation-equivariant and can efficiently capture local interaction motifs such as hydrogen bonds or hydrophobic packing when complexes are encoded as multi-channel atomic density grids.

3.3 Graph neural networks and equivariance

Molecules and proteins are naturally represented as graphs with atoms as nodes and bonds or spatial proximity as edges. Graph neural networks (GNNs) propagate information via message passing:

$$h_i^{(\ell+1)} = \phi(h_i^{(\ell)}, \{m_{ij}^{(\ell)} : j \in \mathcal{N}(i)\}) \quad (6)$$

where messages $m_{ij}^{(\ell)} = \psi(h_i^{(\ell)}, h_j^{(\ell)}, e_{ij})$ depend on neighbour features and edge attributes, such as distances, bond types, and ϕ, ψ are learnable functions. This allows learning local chemical environments and non-covalent interaction patterns. For docking, models must respect Euclidean symmetries: a rigid rotation or translation of the input complex should not change predicted energies, or should predictably transform coordinates. SE(3)-equivariant networks enforce for any rotation R and translation t (**Equation 7**).

$$f(RX + t) = Rf(X) + t \quad (7)$$

Here X denotes the set of atomic coordinates. Such architectures underlie many modern models, including EquiBind[3] DiffDock’s[4] score network, and AlphaFold2’s[5] structure module, enabling more data-efficient learning and predictions consistent with the laws of physics.

4 Deep Learning Scoring Functions

4.1 Machine learning-based scoring

Machine learning-based scoring functions (MLSFs) replace the manually fitted linear combination of descriptors in classical scoring functions by a flexible non-linear mapping learned from data[2]. Early MLSFs such as RF-Score

used Random Forests on simple counts of protein–ligand atom-type pairs within distance cutoffs to predict binding affinity[2, 11]. Later, deep learning approaches introduced 3D CNNs and GNNs operating directly on structural representations.

A generic MLSF can be written as:

$$\hat{\Delta G} = f_{\theta}(\Phi(X)) \quad (8)$$

where $\Phi(X)$ encodes the protein–ligand complex into a feature tensor, like a voxelised grid of atomic densities or a graph, and f_{θ} is a deep network trained on experimental affinities. The loss function is typically the MSE between the predicted and experimental ΔG or pK_d .

4.2 3D CNN scoring

GNINA and CNN rescoring: GNINA[1] represents the protein–ligand complex as a 3D voxel grid with channels for atom types and densities, and trains an ensemble of CNNs to score poses. In its default workflow, GNINA uses a Vina-like empirical scoring function during a Markov chain Monte Carlo (MCMC) search to sample poses, then applies the CNN ensemble to rescore and re-rank the generated poses. This hybrid strategy improves the percentage of targets for which the top-ranked pose lies within 2 Å RMSD of the crystallographic conformation compared to the Vina scoring alone in both redocking and cross-docking benchmarks.

KDEEP and Pafnucy: Methods like KDEEP[12] and Pafnucy[13] similarly encode complexes in 3D grids and use deep CNNs to directly fit binding affinities on PDBBind-like datasets. These models achieve high correlation between predicted and experimental affinities on curated benchmarks. However, their real-world performance depends strongly on data curation and the similarity between training and application domains.

4.3 Graph- and distance-based scoring

An alternative is to build MLSFs on graph or distance-based representations that avoid discretisation artefacts. RTMScore[14] and related approaches, for example, learn residue–atom distance likelihood potentials using mixture density networks, providing a flexible scoring framework that can be tuned for

multiple tasks, such as docking, ranking, and screening[15]. More recent reviews highlight the diversity of architectures, from CNNs to message-passing networks and transformers, and identify problems of bias, overfitting, and evaluation protocols specific to MLSFs[2, 11].

4.4 Caveats

Despite promising benchmarks, MLSFs can exhibit performance degradation under more realistic testing conditions, especially when decoy sets differ from training distributions or when physics-based constraints are not enforced[11]. Recent work emphasises the need for robust evaluation frameworks and diagnostics that simplify genuine learning of physics-based patterns from datasets, highlighting the importance of combining ML with physics-based constraints and post hoc validation such as energy minimisation or MD.[16]

5 Conformation and Pose Generation

5.1 Direct-shot geometric models: EquiBind

EquiBind[3] exemplifies a paradigm shift from sampling-based docking to direct-shot prediction of binding poses using geometric deep learning. Instead of generating and scoring thousands of poses, EquiBind builds SE(3)-equivariant graph representations of the protein and a random ligand conformer, then predicts:

1. The likely binding region, keypoints, on the protein surface.
2. A roto-translation that aligns ligand keypoints to protein keypoints.

The network is trained to minimise the mean-squared error between predicted ligand coordinates and experimentally observed poses, along with auxiliary losses on torsion angles[3, 17]. Because it predicts the final pose in a single forward pass, EquiBind achieves speedups of two orders of magnitude over traditional docking while maintaining competitive RMSD accuracy on many benchmarks.

5.2 Diffusion models for docking: DiffDock

DiffDock[4] frames docking as generative modelling over the manifold of ligand poses, using a diffusion model defined over translation, rotation, and torsion degrees of freedom. In a diffusion model, one defines a forward noising process which gradually destroys structure, and trains a neural network $s_\theta(x_t, t)$ to approximate the score function $\nabla_{x_t} \log p_t(x_t)$ (**Equation 9**). The reverse process iteratively denoises from Gaussian noise back to realistic ligand poses.

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I) \quad (9)$$

For docking, x_t encodes the ligand’s translation, rotation, and torsion angles, while the protein is treated as a conditioning variable via an SE(3)-equivariant graph neural network[4]. Empirically, DiffDock achieves higher success rates (RMSD < 2 Å) on PDBBind than classical docking and earlier deep methods, and maintains notable performance even when docking to computationally predicted (AlphaFold-like) structures[4, 18]. Its diffusion-based generative sampler implicitly explores diverse basins of the binding landscape, sidestepping some combinatorial challenges of explicit search.

5.3 Hybrid and quantum-inspired approaches

Several recent works combine deep generative models with classical or even quantum-inspired optimisation schemes. For example, hybrid workflows use DiffDock or EquiBind to predict candidate poses, followed by local refinement with physics-based energy minimisation or MD, trading a small cost increase for improved accuracy[16, 19]. Quantum-inspired methods encode docking as a combinatorial optimisation problem, which is solved by annealing-like schemes, with learned gradients guiding the search in a latent molecular space[20]. Early results suggest modest improvements over both traditional and purely deep learning-based docking in blind docking scenarios, though these approaches are still nascent[20].

6 Deep Learning for Pose Validation

One criticism of several AI-based docking methods is that they can produce geometrically plausible but physically invalid poses, e.g., with steric clashes

or unrealistic ligand strain, especially when evaluated outside the training set. PoseBusters[16] systematically benchmarked multiple deep docking methods, like DeepDock[21], DiffDock, EquiBind, TankBind[22], and Uni-Mol[23], and traditional tools, such as AutoDock Vina, and GOLD, using additional post-prediction filters for clash detection and force field minimisation. Results showed that while AI methods can achieve competitive RMSD values, they often require explicit physics-based information to ensure realistic geometries. To address this, tools like CompassDock[19] introduce auxiliary modules that estimate ligand strain, clash counts, and interaction quality, enabling post hoc reweighting or fine-tuning of generative models to favour physically and biologically plausible poses.

7 Protein Structure Prediction as a Foundation for Docking

7.1 AlphaFold

AlphaFold (AF)[5] represents a breakthrough in protein structure prediction, achieving near-experimental accuracy for many single-chain proteins. AlphaFold2 (AF2) takes as input an amino acid sequence and derives multiple sequence alignments (MSAs) and template features, which feed into an *Evoformer* block that iteratively refines pairwise residue representations and per-residue embeddings. The structural module then predicts backbone frames and side-chain coordinates, optimising an internal potential that incorporates geometric constraints, such as distance distributions and torsion angles[5, 24]. While AF2 was trained primarily on monomeric and some complex structures, its success made high-quality models available for a large fraction of the proteome, greatly expanding the pool of docking-ready or near-ready protein structures[25, 26]. Nevertheless, AF2 does not explicitly model small-molecule ligands and does not directly output binding affinities, so classical and ML-based docking remain essential for ligand prediction.

AlphaFold3 (AF3)[6] advances beyond AlphaFold2 by predicting the structure and interactions of all biomolecules, including DNA, RNA, and ligands, rather than just proteins, utilising a diffusion-based architecture instead of AlphaFold2’s reliance on MSA-based templates. Key improvements include superior accuracy for complexes, better antibody-antigen prediction, and faster, more versatile modelling of complex, non-protein assemblies.

7.2 Complex and multi-state prediction

Several protocols extend AF2 to complexes, such as AlphaFold-Multimer and community workflows, and to alternative conformational states, e.g., active versus inactive GPCRs[27, 28]. These methods typically adjust MSAs and templates or modify training to encourage specific interaction patterns. For docking, AF2-derived complex or multi-state models can provide more realistic binding sites. However, it is advised to practice caution as AF2 is known to over-stabilise certain conformations or misrepresent flexible or disordered regions crucial for allosteric binding[26, 29, 30].

7.3 Using AF structures in docking workflows

Several studies have shown that AI-predicted protein structures can be successfully used for docking and virtual screening, often with performance comparable to that of X-ray crystal structures when binding sites are well resolved[25, 26]. However, limitations arise for:

- Cryptic or allosteric pockets not visible in the AF model.
- Conformational selection or induced-fit, where the bound state differs markedly from AF’s prediction.
- Cases where low pLDDT (predicted Local Distance Difference Test) regions overlap functional sites.

8 Co-folding Models for Protein–Ligand Complexes

Recent models such as AF3 and various open-source alternatives extend the AF2 paradigm to jointly model proteins, nucleic acids, small molecules and other biomolecules in an all-atom representation[6, 28, 31]. These co-folding models take sequences and/or graphs for multiple chains and ligands, plus optional templates and constraints, and output full 3D complex structures. Internally, they use transformer-like architectures with attention over residues, atoms, and pairwise interactions, often augmented with SE(3)-equivariant geometric reasoning. Conceptually, co-folding replaces the traditional pipeline of *predict receptor structure* $\rightarrow$ *dock ligands* by a joint generative process that

accounts for mutual structural adaptation. In principle, such models can capture induced fit and cryptic pocket opening provided relevant configurations are represented in the training data.

8.1 Boltz-2

Boltz-2[7] is an open-source foundation model designed to co-predict complex 3D structures and binding affinities for diverse biomolecular interactions. Building on earlier models, including Boltz-1[32] and AF-like architectures, Boltz-2:

- Jointly encodes proteins, nucleic acids, and small molecules using multimodal transformers and equivariant geometric modules.
- Predicts 3D complex structures, co-folding, across these modalities.
- Includes dedicated heads for continuous binding affinity and binder or decoy classification.

The training objective combines structural loss terms, such as frame-aligned point error between predicted and experimental coordinates, with affinity losses, e.g., MSE in experimental binding affinities or negative log-likelihood of experimental measurements, effectively learning a surrogate for the binding free energy landscape.

Benchmark studies[7, 33] suggest that Boltz-2 achieves strong correlation with experimental binding affinities and approaches FEP-level accuracy on several small-molecule benchmarks while being approximately three orders of magnitude faster. External evaluations indicate that Boltz-2 is particularly effective on relatively rigid targets and chemotypes well represented in its training set. In contrast, performance can degrade for highly flexible ligands or interaction motifs that are under-represented in the training data.

8.2 Mechanistic perspective

From a physicochemical standpoint, co-folding models, like Boltz-2, can be viewed as learning a parameterised approximation to the joint Boltzmann distribution.

$$p_{\theta}(X, \Delta G \mid S) \approx \frac{1}{Z(S)} \exp \left(-\beta F_{\theta}(X, S) \right) \quad (10)$$

where X denotes the complex coordinates, S the sequences and chemical graphs of the interacting species, ΔG an effective binding free energy, and F_θ a learned free-energy-like functional, $\beta = (\frac{1}{k_b T})$. During inference, the model outputs a mode or sample of this distribution, a predicted complex, and an estimate of ΔG , effectively mapping sequence and ligand structure to both pose and affinity. While F_θ is not a true physical free energy, calibration against experimental data and comparison with FEP suggest that it captures substantial aspects of the underlying free energy landscape[7].

9 Binding Affinity Prediction with Deep Learning

9.1 Supervised affinity regression

Classical MLSFs and new foundation models alike treat binding affinity prediction as supervised regression, with the dataset comprising complexes $X^{(n)}$ and associated experimental labels $y^{(n)}$, such as pK_i or pIC_{50} . A typical loss is

$$\mathcal{L}_{\text{aff}}(\theta) = \frac{1}{N} \sum_{n=1}^N \ell(f_\theta(X^{(n)}), y^{(n)}), \quad (11)$$

where ℓ may be MSE or Huber loss. Some models also incorporate uncertainty estimation, interpreting f_θ as the mean of a conditional distribution with learned variance. An attractive interpretation connects predicted affinity to the Boltzmann distribution of bound v/s unbound states. If one assumes

$$K_d \propto \exp(\beta \Delta G_{\text{bind}}), \quad (12)$$

then learning ΔG_{bind} is equivalent to learning $\log K_d$, which is roughly linear in experimental pK_d . Models like KDEEP[12], Pafnucy[13] and their variants effectively fit such a relationship over large, but imperfect, structural datasets.

9.2 Joint pose and affinity modelling

A key challenge is coupling pose prediction and affinity estimation in a physically meaningful manner. In classical pipelines, docking generates poses, and a separate rescoring or free-energy method estimates affinities, often requiring careful manual curation of poses. New models attempt to couple these steps:

- DiffDock learns a distribution over poses and provides confidence scores that correlate with RMSD, offering a proxy for pose quality and, indirectly, binding likelihood.
- Co-folding models, like Boltz-2[7], directly predict the bound pose and a calibrated affinity from shared internal representations, improving the link between geometric and thermodynamic aspects.
- Hybrid ML/MM pipelines use ML-predicted poses as seeds for physics-based calculations (MM/GBSA, FEP), reducing sampling overhead while maintaining interpretability.

Designing loss functions and training models that maintain consistency between pose and affinity predictions is an active area of research, with approaches ranging from multi-task learning to contrastive or ranking-based objectives that enforce ranking-ordering consistency with experimental SAR data[11, 15].

10 How AF-like and Boltz-2-like Models Predict Complexes?

10.1 Architecture

Although implementation details vary, AF2[5], AF-like co-folding models[24, 28], and Boltz-2[7] share several architectural motifs:

1. *Sequence and pair representations*: Each residue or small molecule fragment is encoded into a high-dimensional embedding using its sequence or graph features. Pairwise residue-residue and residue-ligand relationships are encoded into separate pair representations.

2. *Iterative refinement via attention*: Multi-head attention layers update both per-residue and pairwise features, allowing long-range communication within and across chains, analogous to message passing on a fully connected interaction graph.
3. *Geometric module*: An SE(3)-equivariant structural module converts refined features into 3D frames and atomic coordinates, iteratively adjusting them to satisfy learned geometric constraints.
4. *Additional heads*: Additional heads output per-residue confidence scores, predicted aligned error (PAE) maps, and in Boltz-2 binding affinities and binder or decoy likelihoods.

10.2 Training signals and objectives

For complex prediction, the main geometric loss compares predicted and experimental coordinates via frame-aligned point error or similar metrics:

$$\mathcal{L}_{\text{geom}} = \sum_i \|\hat{\mathbf{r}}_i - \mathbf{r}_i\|^2 \quad (13)$$

After optimal rigid alignment of local frames, additional losses penalise deviations in inter-residue distances, dihedral angles, and clash counts, effectively teaching the model a differentiable surrogate of a physics-based potential[5, 7, 24].

For affinity, Boltz-2 and related models[7] add regression or ranking losses:

$$\mathcal{L}_{\text{aff}} = \frac{1}{N} \sum_n \left(\hat{y}^{(n)} - y^{(n)} \right)^2 \quad (14)$$

optionally combined with contrastive terms that encourage the correct ordering of affinities within the SAR series. Joint training on both structural and affinity data encourages internal representations that are sensitive not only to geometry but also to interaction features correlated with binding affinity.

10.3 Inference

Given a protein sequence and a candidate ligand, an AF3 or Boltz-2[7, 28] like pipeline proceeds broadly as follows:

1. Build MSAs and template features for the protein, featurise the ligand via molecular graph, SMILES, or 3D conformers.
2. Run the encoder, an Evoformer-like block, to produce refined per-residue and cross-modality pair features.
3. Use the geometric head to generate a 3D complex structure, possibly multiple samples via different seeds.
4. Use the affinity head to output a predicted binding affinity and associated confidence or uncertainty.

Since structure and affinity predictions share the same backbone, updates that improve structural accuracy often improve affinity prediction and vice versa. However, careful training and calibration are required to avoid overfitting and handle dataset noise[7, 33].

11 Future Directions and Open Challenges

AI and deep learning have already delivered tangible improvements in docking accuracy, speed, and scope; however, significant challenges remain. Data quality and representativeness in structural and affinity datasets limit the ability of MLSFs and foundation models to generalise beyond well-studied chemotypes and targets. Methods like CompassDock[19], ToolBoxSF[11] and PoseBusters[16] highlight the need for rigorous, physics-informed validation and benchmarking to avoid over-optimistic claims of *state-of-the-art* performance. A particularly interesting direction is integrating generative ligand design, co-folding structure prediction, and high-quality affinity estimation into closed-loop pipelines. Recent studies combining generative deep learning, docking and MD for pocket-aware inhibitor design illustrate the potential of such frameworks, but also underscore the importance of accurate binding models, especially for transient and allosteric sites[30, 34, 35]. Finally, the emergence of quantum-inspired and hardware-accelerated docking, such as FPGA accelerators for DiffDock, suggests that specialised computing stacks will complement algorithmic advances[20, 36].

References

- [1] Andrew T McNutt, Paul Francoeur, Rishal Aggarwal, Tomohide Masuda, Rocco Meli, Matthew Ragoza, Jocelyn Sunseri, and David Ryan Koes. Gnina 1.0: molecular docking with deep learning. *Journal of cheminformatics*, 13(1):43, 2021.
- [2] Rocco Meli, Garrett M Morris, and Philip C Biggin. Scoring functions for protein-ligand binding affinity prediction using structure-based deep learning: a review. *Frontiers in bioinformatics*, 2:885983, 2022.
- [3] Hannes Stärk, Octavian Ganea, Lagnajit Pattanaik, Regina Barzilay, and Tommi Jaakkola. Equibind: Geometric deep learning for drug binding structure prediction. In *International conference on machine learning*, pages 20503–20521, 2022.
- [4] Gabriele Corso, Hannes Stärk, Bowen Jing, Regina Barzilay, and Tommi Jaakkola. Diffdock: Diffusion steps, twists, and turns for molecular docking. *arXiv preprint arXiv:2210.01776*, 2022.
- [5] John Jumper, Richard Evans, Alexander Pritzel, Tim Green, Michael Figurnov, Olaf Ronneberger, Kathryn Tunyasuvunakool, Russ Bates, Augustin Žídek, Anna Potapenko, et al. Highly accurate protein structure prediction with alphafold. *Nature*, 596(7873):583–589, 2021.
- [6] Josh Abramson, Jonas Adler, Jack Dunger, Richard Evans, Tim Green, Alexander Pritzel, Olaf Ronneberger, Lindsay Willmore, Andrew J Ballard, Joshua Bambrick, et al. Accurate structure prediction of biomolecular interactions with alphafold 3. *Nature*, 630(8016):493–500, 2024.
- [7] Saro Passaro, Gabriele Corso, Jeremy Wohlwend, Mateo Reveiz, Stephan Thaler, Vignesh Ram Somnath, Noah Getz, Tally Portnoi, Julien Roy, Hannes Stark, et al. Boltz-2: Towards accurate and efficient binding affinity prediction. *BioRxiv*, 2025.
- [8] Ingrid Daubechies, Ronald DeVore, Simon Foucart, Boris Hanin, and Guergana Petrova. Nonlinear approximation and (deep) relu networks. *Constructive Approximation*, 55(1):127–172, 2022.
- [9] Sridhar Narayan. The generalized sigmoid activation function: Competitive supervised learning. *Information sciences*, 99(1-2):69–82, 1997.

- [10] Fangyu Zou, Li Shen, Zequn Jie, Weizhong Zhang, and Wei Liu. A sufficient condition for convergences of adam and rmsprop. In *Proceedings of the IEEE/CVF Conference on computer vision and pattern recognition*, pages 11127–11135, 2019.
- [11] Guy Durant, Fergus Boyles, Kristian Birchall, Brian Marsden, and Charlotte M Deane. Robustly interrogating machine learning-based scoring functions: what are they learning? *Bioinformatics*, 41(2):btaf040, 2025.
- [12] José Jiménez, Miha Skalic, Gerard Martinez-Rosell, and Gianni De Fabritiis. K deep: protein–ligand absolute binding affinity prediction via 3d-convolutional neural networks. *Journal of chemical information and modeling*, 58(2):287–296, 2018.
- [13] Marta M Stepniewska-Dziubinska, Piotr Zielenkiewicz, and Pawel Siedlecki. Development and evaluation of a deep learning model for protein–ligand binding affinity prediction. *Bioinformatics*, 34(21):3666–3674, 2018.
- [14] Zijie Qiu, Sheng Xu, Junkang Wei, Tao Shen, and Siqi Sun. Rtm-align: an improved rna alignment tool with enhanced short sequence performance via post-standardization and fragment alignment. *bioRxiv*, pages 2024–05, 2024.
- [15] Chao Shen, Xujun Zhang, Chang-Yu Hsieh, Yafeng Deng, Dong Wang, Lei Xu, Jian Wu, Dan Li, Yu Kang, Tingjun Hou, et al. A generalized protein–ligand scoring framework with balanced scoring, docking, ranking and screening powers. *Chemical Science*, 14(30):8129–8146, 2023.
- [16] Martin Buttenschoen, Garrett M Morris, and Charlotte M Deane. Pose-busters: Ai-based docking methods fail to generate physically valid poses or generalise to novel sequences. *Chemical Science*, 15(9):3130–3139, 2024.
- [17] Yuze Li, Li Li, Shuang Wang, and Xiaowen Tang. Equibind: A geometric deep learning-based protein-ligand binding prediction method. *Drug Discoveries & Therapeutics*, 17(5):363–364, 2023.
- [18] Patricia Suriana and Ron O Dror. Enhancing ligand pose sampling for molecular docking. *ArXiv*, pages arXiv–2312, 2023.

- [19] Ahmet Sarigun, Vedran Franke, Bora Uyar, and Altuna Akalin. Compassdock: Comprehensive accurate assessment approach for deep learning-based molecular docking in inference and fine-tuning. *arXiv preprint arXiv:2406.06841*, 2024.
- [20] Runqiu Shu, Bowen Liu, Zhaoping Xiong, Xiaopeng Cui, Yunting Li, Wei Cui, Man-Hong Yung, and Nan Qiao. Quantum-inspired machine learning for molecular docking. *arXiv preprint arXiv:2401.12999*, 2024.
- [21] Zhirui Liao, Ronghui You, Xiaodi Huang, Xiaojun Yao, Tao Huang, and Shanfeng Zhu. Deepdock: enhancing ligand-protein interaction prediction by a combination of ligand and structure information. In *2019 IEEE International Conference on Bioinformatics and Biomedicine (BIBM)*, pages 311–317. IEEE, 2019.
- [22] Wei Lu, Qifeng Wu, Jixian Zhang, Jiahua Rao, Chengtao Li, and Shuangjia Zheng. Tankbind: Trigonometry-aware neural networks for drug-protein binding structure prediction. *Advances in neural information processing systems*, 35:7236–7249, 2022.
- [23] Gengmo Zhou, Zhifeng Gao, Qiankun Ding, Hang Zheng, Hongteng Xu, Zhewei Wei, Linfeng Zhang, and Guolin Ke. Uni-mol: A universal 3d molecular representation learning framework. In *The eleventh international conference on learning representations*, 2023.
- [24] Patrick Cramer. Alphafold2 and the future of structural biology. *Nature structural & molecular biology*, 28(9):704–705, 2021.
- [25] Kathryn Tunyasuvunakool, Jonas Adler, Zachary Wu, Tim Green, Michal Zielinski, Augustin Židek, Alex Bridgland, Andrew Cowie, Clemens Meyer, Agata Laydon, et al. Highly accurate protein structure prediction for the human proteome. *Nature*, 596(7873):590–596, 2021.
- [26] Adeline Goulet and Christian Cambillau. Present impact of alphafold2 revolution on structural biology, and an illustration with the structure prediction of the bacteriophage j-1 host adhesion device. *Frontiers in Molecular Biosciences*, 9:907452, 2022.

- [27] Lim Heo and Michael Feig. Multi-state modeling of g-protein coupled receptors at experimental accuracy. *Proteins: Structure, Function, and Bioinformatics*, 90(11):1873–1885, 2022.
- [28] Andriy Kryshchak, Maciej Antczak, Marta Szachniuk, Tomasz Zok, Rachael C Kretsch, Ramya Rangan, Phillip Pham, Rhiju Das, Xavier Robin, Gabriel Studer, et al. New prediction categories in casp15. *Proteins: Structure, Function, and Bioinformatics*, 91(12):1550–1557, 2023.
- [29] Letícia MF Bertoline, Angélica N Lima, Jose E Krieger, and Samantha K Teixeira. Before and after alphafold2: An overview of protein structure prediction. *Frontiers in bioinformatics*, 3:1120370, 2023.
- [30] Gustav Olanders, Giulia Testa, Alessandro Tibo, Eva Nittinger, and Christian Tyrchan. Challenge for deep learning: protein structure prediction of ligand-induced conformational changes at allosteric and orthosteric sites. *Journal of Chemical Information and Modeling*, 64(22): 8481–8494, 2024.
- [31] Sanggeun Park, Sojung Myung, and Minkyung Baek. Advancing protein structure prediction beyond alphafold2. *Current Opinion in Structural Biology*, 90:102985, 2025.
- [32] Jeremy Wohlwend, Gabriele Corso, Saro Passaro, Noah Getz, Mateo Reveiz, Ken Leidal, Wojtek Swiderski, Liam Atkinson, Tally Portnoi, Itamar Chinn, et al. Boltz-1 democratizing biomolecular interaction modeling. *BioRxiv*, pages 2024–11, 2025.
- [33] Shunzhou Wan, Xibei Zhang, Xiao Xue, and Peter V Coveney. On the reliability of ai methods in drug discovery: Evaluation of boltz-2 for structure and binding affinity prediction. *arXiv preprint arXiv:2603.05532*, 2026.
- [34] Li-Ting Zheng, Kun Qian, Jun Zhang, Meng-Ting Liu, Yi Li, and Li-Quan Yang. Exploring pocket-aware inhibitors of btk kinase by generative deep learning, molecular docking, and molecular dynamics simulations. *RSC advances*, 15(42):35336–35355, 2025.
- [35] Serina L Robinson. Artificial intelligence for microbial biotechnology: beyond the hype. *Microbial Biotechnology*, 15(1):65–69, 2022.

- [36] Chuanzhao Zhang, Rui Liu, Shidi Tang, Shun Li, and Ming Ling. Diffdock-fpga: A hardware accelerator for molecular docking with customized tensor product framework and sparse-aware access strategy. In *Proceedings of the Great Lakes Symposium on VLSI 2025*, pages 435–441, 2025.