Chapter 5: Methods to Identify Binding Sites on Proteins

Elvis Martis

1 Introduction

The identification of binding sites on protein structures is a central problem in structural biology, enzymology, and structure-based drug design. Binding sites are spatially localised regions on the protein surface or within the protein interior where small molecules, peptides, nucleic acids, lipids, or proteins form specific, energetically favourable interactions with selected amino acid residues. Correctly locating these regions is a pre-requisite for molecular docking, virtual screening, fragment-based design, and rational engineering of protein function. This chapter provides a detailed overview of methods for identifying binding sites on proteins, focusing on the conceptual and theoretical foundations of binding pockets, sequence-based and static-structure-based approaches, machine-learning and deep-learning methods, and molecular dynamics-based methods that are instrumental for locating hard-to-find cryptic and transient pockets. We aim to provide a clear, practical methodology for validating predicted sites, especially elusive cryptic or transient pockets.

A minimal definition of a binding site is a set of residues whose side-chains and backbone atoms contribute significantly to the binding free energy of a ligand. From a geometric viewpoint, computational methods often start from the protein's solvent-accessible or solvent-excluded surface and identify cavities as enclosed voids within the protein. Pockets can be defined as surface indentations that can accommodate a ligand. And, grooves or clefts are extended concave regions, often along interfaces. Oftentimes, binding sites are deeply buried within protein structures, making them difficult to locate by visual inspection alone.

Before we delve into the methods to identify binding sites on proteins, we first mention, albeit briefly, the different types of binding sites:

Orthosteric Site: A site where endogenous ligands bind. For example, a substrate binds to the catalytic centre of an enzyme; here, the catalytic site is an orthosteric site.

Allosteric Site: A site located distal from the orthosteric site, but is coupled to function through conformational networks or residue-interaction networks.

Protein–protein interfaces: A site, often shallow and extended patches, that is formed when at least two proteins interact.

Cryptic pockets: Binding cavities that are not apparent in the *apo* structures; however, their formation is linked to ligand- or protein-induced conformational changes.

Transient pockets: Binding cavities that appear and disappear on short timescales along the equilibrium dynamics of the protein; these may or may not be ligand-induced. Some of the cryptic pockets may also transient in nature due to their short timescales in which they open and close.

It remains challenging to locate such sites solely through experimental techniques. In practice, an unusual affinity value or structure-active data warrants deeper exploration with the help of computational methods. Over the last three decades, more than fifty distinct computational methods have been developed for protein–ligand binding site prediction, with a clear evolution from purely geometry-based approaches to modern deep-learning models that integrate sequence, structure, and dynamics information.[1–3]

Since computational methods rely on quantitative descriptors of a pocket, one might define them using the following parameters:

- Volume (V) and surface area (A).
- Depth and exposure, computed as fraction of buried v/s solvent-exposed surface.
- Shape descriptors such as sphericity, elongation, roughness.

- Physico-chemical composition, such as hydrophobic/hydrophilic balance, hydrogen-bond donors/acceptors, charge distribution.

Once a potential site is identified, the next challenge is to differentiate pockets that are amenable to small-molecule binding and modulation of the protein's activity. This leads us to the concept of druggability. Druggability scores typically combine descriptors of size, shape, and physico-chemical complementarity into an empirical function validated against known drug-binding pockets.[3, 4]. For example, many tools adopt a linear model of the form

$$\mathrm{D} = \mathrm{w_V V} + \mathrm{w_A A_{hyd}} + \mathrm{w_P P_{pol}} + \mathrm{w_H H_{bd}} + \mathrm{b} \tag{1}$$

where V is pocket volume, $\mathrm{A_{hyd}}$ hydrophobic surface area, $\mathrm{P_{pol}}$ a polarity or hydrogen-bond capacity descriptor, $\mathrm{H_{bd}}$ counts buried hydrogen-bond donors/acceptors, and the parameters $\mathrm{w_i}$, b are learned from the training set.

2 Thermodynamic aspects of binding sites

Protein–ligand binding is governed by the standard free energy of binding,

$$\Delta \mathrm{G_{bind}} = -\mathrm{k_B T} \ln \mathrm{K_a} = \mathrm{k_B T} \ln \mathrm{K_d} \tag{2}$$

where $\mathrm{K_a}$ and $\mathrm{K_d}$ are the association and dissociation constants, respectively, $\mathrm{k_B}$ is Boltzmann's constant, and T is the absolute temperature. A favourable binding site is one for which specific protein–ligand interactions yield a negative $\Delta\mathrm{G_{bind}}$, arising from enthalpic and entropic contributions. (See the **chapter 2** on scoring function for more detailed discussion)

At the microscopic level, a classical force field description approximates the non-bonded interaction energy between a protein atom i and a ligand atom j as

$$E_{ij} = 4\epsilon_{ij}\left[\left(\frac{\sigma_{ij}}{r_{ij}}\right)^{12} - \left(\frac{\sigma_{ij}}{r_{ij}}\right)^{6}\right] + \frac{q_i q_j}{4\pi\varepsilon_0\varepsilon_r r_{ij}}, \tag{3}$$

where ϵ_{ij} and σ_{ij} are Lennard–Jones repulsive and attractive parameters respectively, $\mathrm{r_{ij}}$ is the interatomic distance, $\mathrm{q_i}$ and $\mathrm{q_j}$ partial charges, and ε_{r} is the effective dielectric constant. Energy-based pocket identification methods exploit such functional forms using molecular probes or small fragments around the protein and compute favourable regions in terms of interaction energy or approximate binding free energy.[4, 5]

2.1 Static v/s dynamic view of the binding sites

A binding site prediction focused on static structures, typically X-ray crystal structures of the *apo* protein or a homolog. Since proteins are inherently dynamic, sampling their conformational ensembles in accordance with the Boltzmann distribution (see **Equation 4**) improves the probability of finding binding sites.

$$p(\mathbf{R}) = \frac{1}{Z} \exp\left[-\beta U(\mathbf{R})\right], \quad \beta = \frac{1}{k_{\mathrm{B}}T}, \tag{4}$$

where R denotes the set of atomic coordinates, U(R) the potential energy, and Z the partition function. Cryptic and transient pockets emerge as low population members of this ensemble: they may not be visible in any single experimental structure. However, they can be captured through sufficiently extensive molecular dynamics simulations using enhanced sampling algorithms[6, 7]. Two schools of thought explain the mechanistic possibilities. One states that the protein transiently samples an open conformation that is favourable for the ligand to enter and bind. While the other states that the ligand induces a conformational change in the protein structure as it approaches the protein. The ligand binding stabilises this new conformation. In practice, both mechanisms can likely coexist, and many computational methods are implicitly tuned to detect pockets that are populated in either regime.

3 Sequence-Based binding site prediction

When only the amino acid sequence is available, or when structural models are uncertain, sequence-based methods provide an efficient approach to binding site prediction. These methods are important for large-scale proteome annotation and for targets lacking high-resolution structures.

3.1 Conservation and motif-based approaches

Functionally important residues, including those that form binding sites, often exhibit greater evolutionary conservation than the surrounding surface. Given a multiple sequence alignment (MSA) of homologous sequences, positional conservation can be quantified using the Shannon entropy (**Equation 5**).

$$H(i) = -\sum_a p_i(a) \log p_i(a), \tag{5}$$

where $p_i(a)$ is the frequency of amino acid type a at an alignment position i. Positions with low entropy (high conservation) are candidates for functional sites. Mapping H(i) or derived conservation scores onto a structure reveals clusters of conserved residues that often coincide with binding sites.[1]

Motif-based methods additionally search for known sequence motifs, such as catalytic triads, metal-binding motifs, or short linear motifs implicated in protein–protein or protein–nucleic acid interactions.[1] While these approaches can suggest regions of functional importance, they lack geometric specificity, and they do not directly predict pocket shape or druggability.

3.2 Machine learning-based methods

Modern sequence-based predictors treat binding site identification as a residue-level classification problem. Each residue i is represented by a feature vector x_i capturing the local sequence context, such as k-mer windows, the evolutionary profiles from position-specific scoring matrices (PSSMs), and the predicted structural features, such as secondary structure, solvent accessibility, and disorder propensities.

A classifier $f(x_i)$ then yields the probability that residue i belongs to a binding site.

For instance, a logistic regression model uses

$$P(y_i = 1 \mid \mathbf{x}_i) = \sigma(\mathbf{w}^\top \mathbf{x}_i + b) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x}_i - b)}, \tag{6}$$

where $y_i \in \{0, 1\}$ indicates non-binding (0) or binding (1) residue, w and b are model parameters, and σ is the logistic function. More powerful models employ random forests, gradient-boosted trees, or deep neural networks trained to minimise cross-entropy loss over labelled residues.

Sequence-based methods are attractive due to their speed and ability to generalise across distant homologs. However, they cannot directly capture 3D pocket geometry or subtle differences between alternative conformational states.[1, 2]

4 Static Structure-based methods

Structure-based binding site identification assumes that a 3D structure of the target or a close homolog is known either by experimental methods or computed structure models. These methods may be classified into:

- Geometry-based approaches.
- Energy-based (probe/fragments) approaches.
- Template- or motif-based structure comparison.
- Hybrid and consensus methods.

4.1 Geometry-Based Pocket Detection

4.1.1 Surface and Grid Representations

The most common starting point is the solvent-excluded surface (SES) or Connolly surface, obtained by rolling a probe sphere, typically radius 1.4 Å over the van der Waals surface of the protein. Cavities and pockets correspond to concave regions of this surface. Alternatively, methods discretise space into a 3D grid and classify grid points as lying inside the protein, in solvent, or in putative cavities[2, 3].

A simple grid-based method represents the protein by a binary occupancy function O(r) (**Equation 7**) defined on the grid points $\mathrm{r_k}$.

$$O(\mathbf{r}_k) = \begin{cases} 1, & \text{if } \min_j \|\mathbf{r}_k - \mathbf{R}_j\| < r_{\text{vdW},j}, \\ 0, & \text{otherwise}, \end{cases} \tag{7}$$

where $\mathrm{R_j}$ and $\mathrm{r_{vdW,j}}$ are the coordinates and van der Waals radius of atom j. Pocket points are identified as solvent-exposed grid points adjacent to the protein, but enclosed from the bulk solvent.[8]

4.1.2 Alpha Shapes, Voronoi Diagrams, and Fpocket

More sophisticated geometric approaches employ alpha shapes[9] and Voronoi diagrams[10]. In brief, Voronoi tessellation partitions the space into cells around each atom; the dual graph is the Delaunay triangulation. An *alpha complex* is then constructed by retaining only simplices, such as vertices, edges, and tetrahedra, with circumspheres of radius less than a chosen parameter α. Cavities correspond to voids in this alpha complex.[3]

Fpocket[3] is another widely used geometry-based method that detects pockets by clustering so-called *alpha spheres*. These spheres are tangent to four protein atoms, and contain no other atomic centers. Fpocket search begins with computing an initial set of alpha spheres from a Voronoi tessellation. Followed by filtering the alpha spheres based on radius to select those likely to represent pocket-like voids. Next, it generates clusters of nearby alpha spheres into pockets using a distance-based clustering method. Lastly, it calculates pocket descriptors, such as volume, polarity, and hydrophobicity, and assigns a druggability score using a linear model similar to Equation 1.

4.1.3 Scale Space and Difference of Gaussians (DoGSite)

DoGSite[11, 12] constructs a 3D occupancy grid for the protein, and then applies Gaussian smoothing over multiple length scales. Regions that persist as density minima across scales are designated as pockets. If $\rho(\mathrm{r})$ is a smooth occupancy field, the difference of Gaussians operator is given by equation 8.

$$L(\mathbf{r};\sigma_1,\sigma_2) = G(\mathbf{r};\sigma_1) * \rho(\mathbf{r}) - G(\mathbf{r};\sigma_2) * \rho(\mathbf{r}), \tag{8}$$

where G is a Gaussian kernel of width σ and * denotes convolution, which highlights features on a characteristic length scale. Persistent negative regions of L correspond to concavities. This multiscale view makes DoGSite robust to noise and able to detect both small and large pockets.[2, 11]

4.2 Energy-based and Probe mapping methods

Geometry alone does not guarantee druggability; some pockets may be geometrically prominent yet energetically unfavourable for ligand binding, e.g., highly polar, or fully solvent-exposed grooves. Energy-based approaches address this by mapping interaction energies of small probe molecules onto the protein surface.

4.2.1 FTSite and FTMap

FTSite[5] is an energy-based binding-site identification method built on the FTMap fragment-mapping algorithm[13]. The basic steps are:

1. A set of diverse small organic probe molecules, representing different chemotypes, is docked or sampled exhaustively over the protein surface, using a rigid-body search and fast Fourier transform (FFT) score.

2. For each probe type, low-energy poses are clustered to identify consensus sites.

3. Consensus clusters across different probe types are combined to yield binding hot spots. These clusters of hot spots define putative ligand binding sites.

In essence, FTSite[5] approximates the probability density of finding favourable protein-ligand interactions as given by equation 9.

$$p(\mathbf{r}) \propto \sum_{m} \sum_{k \in \mathcal{C}_m} \exp\left[-\beta E_{mk}\right] \delta(\mathbf{r} - \mathbf{r}_{mk}), \tag{9}$$

where m indexes probe types, k configurations within cluster C_m, and E_{mk} is the interaction energy as in Equation 3. Regions with high p(r) across diverse probe types suggest druggable sites.

4.2.2 Grid-based interaction fields

An alternative is to compute grid-based interaction fields (IFPs or ΔG grids) with one or more probe types, such as hydrophobic carbon, hydrogen-bond donor/acceptor, cation, or anion. At each grid point r_k, one evaluates an approximate interaction free energy $\Delta G_t(r_k)$ for probe type t. Sites with multiple grid points where ΔG_t is strongly favourable across complementary probe types are identified as binding hot spots. Such approaches underlie commercial tools like SiteMap[4].

4.3 Template- and Motif-Based Methods

Template-based methods[1] exploit the observation that binding sites tend to be more conserved within a structure than across overall folds. Given a query protein, these algorithms:

1. Maintain a library of known binding sites described as local structural motifs, e.g., sets of residues with specific geometric and chemical patterns.

2. Superpose these motifs onto the query structure using local or global alignment of backbone or side-chain atoms.

3. Score matches by structural RMSD and chemical similarity, transferring binding site annotations from templates to the query.

FINDSITE[8], for example, uses threading to identify distant homologs in the PDB and then transfers binding site information based on conserved ligand-binding residues across structurally aligned templates. SiteMotif[14] uses graph-based pattern recognition to derive recurrent structural motifs of ligand binding sites and then searches for their occurrences in target proteins.

Template-based methods[1] can detect subtle pockets that would be difficult to identify solely from geometry, particularly for cofactors or metal-binding sites. Their limitations include dependence on the coverage and diversity of template libraries and potential bias toward well-studied folds.

5 Machine dearning and deep learning methods

5.1 Machine learning on geometric cavities

A natural progression from purely geometric methods is to combine geometric cavity detection with supervised learning to distinguish true ligand-binding pockets from other surface concavities. Methods such as P2Rank[15] and related tools follow this paradigm[2].

A typical workflow comprises:

1. Detect all candidate cavities using a geometric method, e.g., Fpocket, DoGSite, custom grid-based algorithms.

2. Compute cavity level descriptors: volume, depth, hydrophobicity, electrostatics, residue composition, backbone rigidity, secondary structure, and conservation scores.

3. Train a classifier on labelled data of binding v/s non-binding cavities to predict a probability of being a ligand-binding site.

P2Rank[15], for instance, uses a point-based representation of the protein surface and learns to assign a druggability score to surface points, which are then clustered into pockets. Such ML-based scoring methods significantly improve the ranking of true binding sites compared to purely physics- or rule-based scoring[1, 2].

5.2 Deep Learning on 3D Grids and Surfaces

Deep learning methods treat binding site prediction as a segmentation or detection problem on 3D molecular data. Representative architectures include 3D convolutional neural networks (3D-CNNs), 3D U-Nets, and point cloud networks[16].

5.2.1 3D Convolutional Neural Networks

In 3D-CNN approaches, the protein is embedded in a voxel grid. Each voxel contains multiple channels encoding local atomic densities or properties, such as atom type, partial charge, and hydrophobicity. A 3D convolutional layer computes

$$y_{c'}(\mathbf{r}) = \sigma\left(\sum_{c}\sum_{\boldsymbol{\delta}} W_{c',c}(\boldsymbol{\delta})\, x_c(\mathbf{r}+\boldsymbol{\delta}) + b_{c'}\right) \tag{10}$$

where x_c and $y_{c'}$ are input and output feature maps, $W_{c',c}$ is a learnable 3D kernel, δ indexes positions within the kernel, and σ is a non-linear activation function. Stacks of such layers capture increasingly global spatial patterns, and the network is trained to output voxel-level probabilities of being inside a binding site[16]. Methods like DeepSite[17] and Kalasanty[18] employ this strategy, yielding high-resolution pocket segmentation and druggability estimates.

5.2.2 Surface and Point-Cloud Networks

Protein surfaces can be represented as point clouds[19] or meshes, each with associated features such as curvature, electrostatic potential, and residue identity. Graph-based or point-based networks, such as PointNet[20] and disordered graph convolutional neural network (DGCNN)[21] variants, can operate directly on these representations. For instance, point-based architectures convert atoms into a point cloud and apply U-Net-like architectures[22] with sparse 3D convolutions to segment binding regions on protein–protein interfaces.

5.3 Graph Neural Networks and Geometric Deep Learning

Graph neural networks (GNNs) represent proteins as graphs with nodes corresponding to residues or atoms and edges encoding spatial proximity or chemical bonds[16, 23]. A generic message-passing layer in a GNN can be written as,

$$\mathbf{h}_i^{(l+1)} = \phi\left(\mathbf{h}_i^{(l)}, \sum_{j\in\mathcal{N}(i)} \psi\left(\mathbf{h}_i^{(l)}, \mathbf{h}_j^{(l)}, \mathbf{e}_{ij}\right)\right), \tag{11}$$

where $\mathrm{h}_i^{(l)}$ is the feature vector of node i at layer l, N(i) is the neighborhood of i, e_{ij} edge features, such as distance, orientation, and ϕ, ψ are neural networks. After several layers, node-level outputs can be interpreted as binding propensities for residues or atoms. GNNs are particularly appealing for cryptic and allosteric site prediction when combined with dynamic information. PocketMiner[23], for example, is a graph neural network trained to predict residues likely to be part of the cryptic pockets that open in MD simulations; it is reported to be orders of magnitude faster than MD-based screening while retaining good accuracy.

5.4 Deep Learning on Sequences and Structures

Advances in protein language models, such as ESM[24] and ProteinBert[25], and in structure prediction methods like AlphaFold and RoseTTAFold have enabled hybrid models that combine sequence embeddings and structural features. A typical architecture might:

- Encode the sequence using a transformer-based language model to obtain residue embeddings s_i.
- Encode the 3D structure or predicted distance maps into structural embeddings t_i.
- Concatenate or fuse s_i and t_i through attention layers to predict residue-level binding probabilities.

The attention mechanism computes

$$\text{Attention}(\mathbf{Q},\mathbf{K},\mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_k}}\right)\mathbf{V}, \tag{12}$$

where queries Q, keys K, and values V are linear projections of $\{s_i, t_i\}$, and d_k is the key dimension. Such models can naturally learn long-range couplings between residues, which are essential for recognising distributed allosteric networks and extended binding sites.

6 Dynamics-based methods

Static approaches can only capture pockets present in a single conformation. Therefore, to reveal cryptic and transient pockets, it is necessary to sample protein dynamics[6, 7].

6.1 Classical molecular dynamics and pocket ensembles

Standard MD simulations generate time series of atomic coordinates $\{R(t)\}$ by integrating Newton's equations of motion under a molecular mechanics force field (see **chapter 9** on MD simulations for detailed explanation). Pocket analysis tools such as MDpocket[26] or TRAPP[27] can detect cavities in each snapshot and accumulate statistics over the trajectory.

Let P_s denote the set of grid points assigned to pocket p in snapshot s. The occupancy of grid point k by pocket p over a trajectory of S snapshots is

$$O_p(k) = \frac{1}{S}\sum_{s=1}^{S} \mathbf{1}\{k \in \mathcal{P}_s\}, \tag{13}$$

where $1\{\cdot\}$ is the indicator function. High values of $O_{p(k)}$ indicate persistent pockets, while intermediate values reflect transient pockets that open and close during the simulation. Pocket volume distributions over time, $P_{p(V)}$, provide insights into the dynamics of pocket opening. For cryptic pockets, one often observes a low baseline volume punctuated by rare opening events with significantly larger volumes[6].

6.2 Cosolvent and Mixed-Solvent MD

Mixed-solvent MD (MSMD), also called cosolvent MD, enhances sampling of binding hot spots by adding small organic probes, like isopropanol, acetonitrile, and benzene, at low molar fractions to the aqueous solvent.[6, 7] Probes preferentially accumulate in regions of the protein that can favourably accommodate hydrophobic or polar groups, thereby highlighting potential binding sites.

The probe density for species t at grid point $\mathrm{r_k}$ is estimated as given in equation 14.

$$\rho_t(\mathbf{r}_k) = \frac{1}{S} \sum_{s=1}^{S} \sum_{m \in \mathcal{M}_t(s)} \mathbf{1}\{\|\mathbf{r}_k - \mathbf{r}_m(s)\| < r_{\mathrm{cut}}\}, \tag{14}$$

where $\mathrm{M_{t(s)}}$ is the set of probe molecules of type t at snapshot s, $\mathrm{r_{m(s)}}$ their coordinates, and $\mathrm{r_{cut}}$ a cutoff distance. Hot spots are defined as regions where $\rho_\mathrm{t}(\mathrm{r_k})$ exceeds a threshold for one or more probe types. MSMD provides both the location of hot spots and dynamic information such as occupation times and opening frequencies. However, cryptic pockets that require large backbone movements may still be considered rare events even with cosolvents. Such rare event sampling warrants the use of enhanced sampling and recent developments such as CrypToth, which integrate MSMD with topological data analysis[7, 28]

6.3 Enhanced Sampling Methods

Enhanced sampling techniques aim to accelerate sampling of conformational space beyond timescales accessible to classical MD simulations. For cryptic pocket discovery, rare events are transitions between closed and open pocket conformations that may require large backbone movements or the rotation of buried residues with significantly large barriers. Methods include accelerated MD, metadynamics, Hamiltonian scaling, e.g., SWISH[29], SWISH-X[30], and weighted-ensemble MD[6, 31].

6.3.1 Accelerated MD and Ligand-Mapping MD

Accelerated MD (aMD) modifies the potential energy surface by adding a bias potential that raises energy wells below a threshold E,

$$U'(\mathbf{R}) = \begin{cases} U(\mathbf{R}) + \Delta U(\mathbf{R}), & U(\mathbf{R}) < E, \\ U(\mathbf{R}), & U(\mathbf{R}) \geq E, \end{cases} \tag{15}$$

with

$$\Delta U(\mathbf{R}) = \frac{(E - U(\mathbf{R}))^2}{\alpha + (E - U(\mathbf{R}))}, \tag{16}$$

where α controls the degree of acceleration, whihc flattens energy wells, enhancing barrier crossing and conformational transitions. Ligand-mapping MD (LMMD)[32] combines MD with small probe ligands to identify binding hot spots dynamically. Accelerated ligand-mapping MD (aLMMD)[33] merges aMD with LMMD to efficiently detect deeply buried cryptic pockets and occluded binding sites.

6.3.2 SWISH / SWISH-X

The SWISH (Sampling Water Interfaces through Scaled Hamiltonians)[29] protocol scales protein–water interactions to facilitate opening of hydrophobic pockets by effectively making water less favourable at certain interfaces. In SWISH-X[30], SWISH is combined with On-the-fly Probability Enhanced Sampling (OPES) MultiThermal to improve sampling across a spectrum of effective temperatures. These methods explore an ensemble of modified Hamiltonians,

$$U_\lambda(\mathbf{R}) = U_{\text{protein}}(\mathbf{R}) + \lambda U_{\text{protein}-\text{water}}(\mathbf{R}) + U_{\text{water}}(\mathbf{R}), \tag{17}$$

with $\lambda < 1$ temporarily weakening protein–water interactions to encourage dehydration and pocket opening; proper reweighting recovers observables of the original system. SWISH-X[30] has been shown to outperform standard MSMD and SWISH[29] in opening cryptic pockets at protein–protein interfaces and isolated proteins, revealing cavities inaccessible to unbiased simulations in similar computational time.

6.3.3 Weighted-Ensemble MD for Pocket Opening

Weighted-ensemble (WE) MD[34] methods focus computational effort on rare transitions by partitioning the CV space (e.g., pocket volume) into bins and spawning or pruning trajectories based on their weights. This enables direct estimation of transition rates between closed and open pocket states and facilitates mapping of pocket opening pathways.

6.3.4 Markov State Models

Extended MD simulations analysed with Markov state models (MSMs)[35] have been used to dissect the mechanisms and kinetics of cryptic pocket opening. In an MSM, the conformational space is discretised into states, such as microstates, based on structural clustering, and a transition probability matrix $\mathrm{T}(\tau)$ is estimated for lag time τ,

$$T_{ij}(\tau) = P(\text{state } j \text{ at } t+\tau \mid \text{state } i \text{ at } t). \tag{18}$$

Diagonalisation of T yields implied timescales for slow processes, including pocket opening and closing. MSMs can reveal which conformational pathways lead to cryptic pocket formation and identify gating residues whose motions control access[6].

6.3.5 Data-Centric and Exposon-Based Methods

Data-centric approaches leverage large MD datasets and unsupervised learning to identify correlated solvent exposure changes, exposons, that signal cryptic site formation[36]. Here, per-residue solvent-accessible surface area (SASA) time series $\mathrm{A}_{\mathrm{i(t)}}$ are clustered into groups of residues with correlated exposure dynamics; clusters that transiently increase SASA together often delineate cryptic pockets.

Formally, one may compute a correlation matrix C_{ij} of SASA fluctuations,

$$C_{ij} = \frac{\langle \delta A_i(t)\, \delta A_j(t) \rangle}{\sqrt{\langle \delta A_i(t)^2 \rangle \langle \delta A_j(t)^2 \rangle}}, \tag{19}$$

cluster residues based on C_{ij}, and then analyse structural mapping of these clusters to identify putative cryptic pockets[6, 36].

6.3.6 CrypToth: MSMD with Topological Data Analysis

CrypToth[28] is a recent method that couples MSMD with topological data analysis (TDA), specifically persistent homology, to distinguish cryptic sites from generic hot spots. After running MSMD with multiple probe types, probe occupancy maps are converted into scalar fields on the protein surface. Hot spots whose topological signatures persist over a broad range of thresholds and correlate with conformational flexibility are ranked highly as cryptic pocket candidates[7, 28]. CrypToth[28] integrates the lifetimes of pocket opening and closing with energetic and dynamic metrics to score sites.

7 Validation of Predicted Binding Sites

Prediction alone is insufficient; computationally identified sites must be validated and prioritised for use as plausible molecular docking sites. Validation strategies include, internal consistency checks, benchmarking against known datasets, energetic analyses, evolutionary evidence, and direct experimental measurements.

7.1 Benchmarking and Performance Metrics

Comparative evaluations of binding site predictors have standardised performance metrics and benchmark datasets[1, 2]. Given a protein with one or more known binding sites and a predictor that outputs a ranked list of candidate pockets, common metrics include:

- Top N recall: A fraction of proteins for which at least one true binding site overlaps with the top N predicted sites.
- Precision and recall at residue or grid-point level:

  $$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN}, \tag{20}$$

  where TP are true positives (correctly predicted binding residues), FP are false positives, and FN are false negatives.
- F1-score:

  $$F_1 = 2\frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}. \tag{21}$$

7.2 Cross-Method Consistency and Consensus

A validation step is to compare results across multiple methods:

- Agreement between geometry-based tools (Fpocket, DoGSite), energy-based mapping (FTSite, MSMD), and ML predictors (P2Rank, DL-based) supports the significance of a site.

- High cryptic pocket scores from PocketMiner or CryptoSite that coincide with pockets seen in enhanced MD or MSMD simulations improve confidence.

Consensus strategies can be formalised through ensemble averaging or voting schemes. For example, given method specific scores $s_{m(p)}$ for pocket p, a consensus score might be

$$S_{\text{cons}}(p) = \sum_m \alpha_m \, \hat{s}_m(p), \tag{22}$$

where $\hat{s}_m$ are normalized scores and α_m weights chosen based on prior performance or cross-validation.

7.3 Druggability and Energetic Validation

Druggability assessment combines geometric and energetic information to estimate whether a pocket is likely to bind small molecules with high affinity[3, 4]. Beyond empirical scoring (**Equation 1**), one can perform:

- Docking of fragment libraries into the predicted site and analysing enrichment of high-scoring poses.
- Binding free energy estimates for representative ligands using end-point methods such as MM/GBSA:

  $$\Delta G_{\text{bind}}^{\text{MM/GBSA}} = \langle E_{\text{complex}} \rangle - \langle E_{\text{protein}} \rangle - \langle E_{\text{ligand}} \rangle + \Delta G_{\text{solv}} - T\Delta S, \tag{23}$$

  where $\langle E \rangle$ are average gas-phase energies and ΔG_{solv} solvation contributions, often from a generalised Born model.
- Alchemical free energy calculations for ligand binding to different pockets or different conformational states (e.g., open v/s closed cryptic pocket) to quantify relative affinities.

For cryptic or transient pockets, a key validation is whether ligands designed or docked into the open pocket remain bound and stable in unbiased MD simulations initiated from those complexes, indicating a realistic environment[6, 36]

7.4 Evolutionary and Coevolutionary Signals

Mapping evolutionary conservation and co-evolutionary couplings onto predicted pockets provides orthogonal evidence of functional relevance[1]. Highly conserved residues located within a predicted pocket suggest functional or regulatory roles. Co-evolutionary analysis (e.g., direct coupling analysis) can identify networks of residues that co-vary, many of which participate in allosteric pathways connecting distant sites to functional regions. For cryptic pockets, conservation may be weaker, as these sites often confer specificity or allosteric regulation unique to subfamilies. Nonetheless, clusters of moderately conserved residues near cryptic pockets can signal evolutionary tuning of allosteric control[7, 31].

8 Practical Workflows for Binding Site Identification

8.1 Static Structure-based workflow

For a newly determined structure (experimental or high-quality computed structure model):

1. Preprocessing: Assign protonation states, add missing side-chains or loops, and relax the structure with restrained MD or minimisation.

2. Geometry-based screening: Run Fpocket, DoGSite, and related tools to identify surface pockets and obtain initial druggability scores.

3. Energy-based mapping: Apply FTMap/FTSite or equivalent probe mapping to confirm energetically favourable hot spots within geometric pockets.

4. ML-based scoring: Use P2Rank or similar ML predictors to re-rank pockets by druggability.

5. Template search: Compare against binding-site templates, such as Site-Motif, FINDSITE, to detect conserved motifs or cofactor sites.

6. Consensus and prioritisation: Integrate scores into a consensus ranking; select top candidates for docking or experimental follow-up.

8.2 Dynamics-based workflow

When cryptic or transient pockets are suspected in protein–protein interfaces, previously classified as undruggable targets, or kinases with known allosteric regulation:

1. Fast pre-screening: Run PocketMiner or related ML tools to identify residues with high cryptic-pocket propensity.

2. Classical MD: Perform multiple independent MD simulations and analyse pocket dynamics using MDpocket or similar tools, computing occupancy and volume distributions (Equation 13).

3. MSMD: Run mixed-solvent MD with diverse molecular probes to map hot spots and correlate with dynamically opening pockets.

4. Enhanced sampling: For rarely sampled pockets, apply SWISH-X, aLMMD, or WE MD to enhance sampling of opening events and estimate kinetics.

5. Post hoc ML/TDA analysis: Apply methods like CrypToth to rank MSMD hot spots and quantify cryptic pocket robustness using persistent homology and dynamic flexibility metrics.

6. Ligand design and validation: design fragments or small molecules targeting the cryptic pocket, validate using molecular docking, MD stability, and experimental assays.

References

[1] Dario Ghersi and Roberto Sanchez. Beyond structural genomics: computational approaches for the identification of ligand binding sites in protein structures. *Journal of structural and functional genomics*, 12 (2):109–117, 2011.

[2] Jingtian Zhao, Yang Cao, and Le Zhang. Exploring the computational methods for protein-ligand binding site prediction. *Computational and structural biotechnology journal*, 18:417–426, 2020.

[3] Vincent Le Guilloux, Peter Schmidtke, and Pierre Tuffery. Fpocket: an open source platform for ligand pocket detection. *BMC bioinformatics*, 10(1):168, 2009.

[4] Thomas A Halgren. Identifying and characterizing binding sites and assessing druggability. *Journal of chemical information and modeling*, 49(2):377–389, 2009.

[5] Chi-Ho Ngan, David R Hall, Brandon Zerbe, Laurie E Grove, Dima Kozakov, and Sandor Vajda. Ftsite: high accuracy detection of ligand binding sites on unbound protein structures. *Bioinformatics*, 28(2):286–287, 2012.

[6] Antonija Kuzmanic, Gregory R Bowman, Jordi Juarez-Jimenez, Julien Michel, and Francesco L Gervasio. Investigating cryptic binding sites by molecular dynamics simulations. *Accounts of chemical research*, 53(3): 654–661, 2020.

[7] Dorota Gašparíková, Rupesh Chikhale, Jason Cole, and Ehmke Pohl. Recent computational advances in the identification of cryptic binding sites for drug discovery. *Bioinformatics Advances*, 5(1):vbaf156, 2025.

[8] Michal Brylinski and Jeffrey Skolnick. A threading-based method (findsite) for ligand-binding site prediction and functional annotation. *Proceedings of the National Academy of sciences*, 105(1):129–134, 2008.

[9] Kristin Tøndel, Endre Anderssen, and Finn Drabløs. Protein alpha shape similarity analysis (passa): a new method for mapping protein binding sites. application in the design of a selective inhibitor of tyrosine kinase 2. *Journal of computer-aided molecular design*, 16(11):831–840, 2002.

[10] Gordon M Crippen. Voronoi binding site models. *Journal of computational chemistry*, 8(7):943–955, 1987.

[11] Andrea Volkamer, Axel Griewel, Thomas Grombacher, and Matthias Rarey. Analyzing the topology of active sites: on the prediction of pockets and subpockets. *Journal of chemical information and modeling*, 50(11):2041–2052, 2010.

[12] Andrea Volkamer, Daniel Kuhn, Friedrich Rippmann, and Matthias Rarey. Dogsitescorer: a web server for automatic binding site prediction, analysis and druggability assessment. *Bioinformatics*, 28(15):2074–2075, 2012.

[13] Chi Ho Ngan, Tanggis Bohnuud, Scott E Mottarella, Dmitri Beglov, Elizabeth A Villar, David R Hall, Dima Kozakov, and Sandor Vajda. Ftmap: extended protein mapping with user-selected probe molecules. *Nucleic acids research*, 40(W1):W271–W275, 2012.

[14] Santhosh Sankar and Nagasuma Chandra. Sitemotif: A graph-based algorithm for deriving structural motifs in protein ligand binding sites. *PLoS Computational Biology*, 18(2):e1009901, 2022.

[15] Radoslav Krivák and David Hoksza. P2rank: machine learning based tool for rapid and accurate prediction of ligand binding sites from protein structure. *Journal of cheminformatics*, 10(1):39, 2018.

[16] Minjie Mou, Zhichao Zhang, Ziqi Pan, and Feng Zhu. Deep learning for predicting biomolecular binding sites of proteins. *Research*, 8:0615, 2025.

[17] José Jiménez, Stefan Doerr, Gerard Martínez-Rosell, Alexander S Rose, and Gianni De Fabritiis. Deepsite: protein-binding site predictor using 3d-convolutional neural networks. *Bioinformatics*, 33(19):3036–3042, 2017.

[18] Marta M Stepniewska-Dziubinska, Piotr Zielenkiewicz, and Pawel Siedlecki. Improving detection of protein-ligand binding sites with 3d segmentation. *Scientific reports*, 10(1):5035, 2020.

[19] Yeji Wang, Shuo Wu, Yanwen Duan, and Yong Huang. A point cloud-based deep learning strategy for protein–ligand binding affinity prediction. *Briefings in bioinformatics*, 23(1):bbab474, 2022.

[20] Charles R Qi, Hao Su, Kaichun Mo, and Leonidas J Guibas. Pointnet: Deep learning on point sets for 3d classification and segmentation. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 652–660, 2017.

[21] Bo Wu, Yang Liu, Bo Lang, and Lei Huang. Dgcnn: Disordered graph convolutional neural network based on the gaussian mixture model. *Neurocomputing*, 321:346–356, 2018.

[22] Yongchang Liu, Peiying Li, Shikui Tu, and Lei Xu. Refinepocket: An attention-enhanced and mask-guided deep learning approach for protein binding site prediction. *IEEE/ACM Transactions on Computational Biology and Bioinformatics*, 20(5):3314–3321, 2023.

[23] Artur Meller, Michael D Ward, Jonathan H Borowsky, Jeffrey M Lotthammer, Meghana Kshirsagar, Felipe Oviedo, Juan Lavista Ferres, and Gregory Bowman. Predicting the locations of cryptic pockets from single protein structures using the pocketminer graph neural network. *Biophysical journal*, 122(3):445a, 2023.

[24] Zeming Lin, Halil Akin, Roshan Rao, Brian Hie, Zhongkai Zhu, Wenting Lu, Nikita Smetanin, Robert Verkuil, Ori Kabeli, Yaniv Shmueli, et al. Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science*, 379(6637):1123–1130, 2023.

[25] Nadav Brandes, Dan Ofer, Yam Peleg, Nadav Rappoport, and Michal Linial. Proteinbert: a universal deep-learning model of protein sequence and function. *Bioinformatics*, 38(8):2102–2110, 2022.

[26] Peter Schmidtke, Axel Bidon-Chanal, F Javier Luque, and Xavier Barril. Mdpocket: open-source cavity detection and characterization on molecular dynamics trajectories. *Bioinformatics*, 27(23):3276–3285, 2011.

[27] Daria B KOKH, Stefan RICHTER, Stefan HENRICH, Paul CZODROWSKI, Friedrich RIPPMANN, and Rebecca C WADE. Trapp: A tool for analysis of transient binding pockets in proteins. *Journal of chemical information and modeling*, 53(5):1235–1252, 2013.

[28] Jun Koseki, Chie Motono, Keisuke Yanagisawa, Genki Kudo, Ryunosuke Yoshino, Takatsugu Hirokawa, and Kenichiro Imai. Cryptoth: Cryptic pocket detection through mixed-solvent molecular dynamics simulations-based topological data analysis. *Journal of Chemical Information and Modeling*, 65(11):5567–5575, 2025.

[29] Federico Comitani and Francesco Luigi Gervasio. Exploring cryptic pockets formation in targets of pharmaceutical interest with swish. *Journal of chemical theory and computation*, 14(6):3321–3331, 2018.

[30] Alberto Borsatto, Eleonora Gianquinto, Valerio Rizzi, and Francesco Luigi Gervasio. Swish-x, an expanded approach to detect cryptic pockets in proteins and at protein–protein interfaces. *Journal of Chemical Theory and Computation*, 20(8):3335–3348, 2024.

[31] Si Zhang and Gregory R Bowman. Decrypting cryptic pockets with physics-based simulations and artificial intelligence. *Current Opinion in Structural Biology*, 96:103215, 2026.

[32] Justin Tze-Yang Ng and Yaw Sing Tan. Accelerated ligand-mapping molecular dynamics simulations for the detection of recalcitrant cryptic pockets and occluded binding sites. *Journal of Chemical Theory and Computation*, 18(3):1969–1981, 2022.

[33] Justin Tze-Yang Ng and Yaw Sing Tan. Accelerated ligand-mapping molecular dynamics simulations for the detection of recalcitrant cryptic pockets and occluded binding sites. *Journal of Chemical Theory and Computation*, 18(3):1969–1981, 2022.

[34] Daniel M Zuckerman and Lillian T Chong. Weighted ensemble simulation: review of methodology, applications, and software. *Annual review of biophysics*, 46:43–57, 2017.

[35] Brooke E Husic and Vijay S Pande. Markov state models: From an art to a science. *Journal of the American Chemical Society*, 140(7): 2386–2396, 2018.

[36] Rommie E Amaro. Will the real cryptic pocket please stand out? *Biophysical Journal*, 116(5):753–754, 2019.