

Methods to Identify Binding Sites on Proteins

Elvis Martis

1 Introduction

The identification of ligand binding sites on protein structures is a central problem in structural biology, enzymology, and structure-based drug design. Binding sites constitute spatially localized regions on the protein surface or in the protein interior where ligands—small molecules, peptides, nucleic acids, lipids, or proteins—form specific and energetically favorable interactions with selected amino acid residues. Correctly locating these regions is a prerequisite for molecular docking, virtual screening, fragment-based design, and rational engineering of protein function.

Classically, binding sites have been associated with obvious surface clefts or cavities, for example catalytic pockets in enzymes. However, the modern view is substantially richer. Binding sites may be:

- *Orthosteric*, where endogenous ligands bind.
- *Allosteric*, located distal from the orthosteric site but coupled to function via conformational networks.
- *Protein–protein interfaces*, which often form shallow and extended patches.
- *Cryptic pockets*, cavities that are not apparent in ground-state structures but emerge upon conformational excursions or ligand binding.
- *Transient pockets*, cavities that appear and disappear on short timescales along the equilibrium dynamics of the protein.

High-throughput experimental discovery of such sites remains challenging. In contrast, computational methods can analyze large numbers of protein structures or models with modest cost, providing hypotheses for downstream experimental validation. Over the last three decades, more than fifty distinct computational methods have been developed for protein–ligand binding site prediction, with a clear evolution from purely geometry-based

approaches to sophisticated machine-learning and deep-learning models that integrate sequence, structure, and dynamics information.[5, 18, 10]

This chapter provides a detailed overview of methods to identify binding sites on proteins, focusing on:

1. Conceptual and theoretical foundations of binding pockets.
2. Sequence-based and static structure-based approaches.
3. Machine-learning and deep-learning methods.
4. Dynamics-based and enhanced-sampling strategies, with emphasis on cryptic and transient pockets.
5. Practical strategies for validating predicted sites, especially elusive cryptic or transient pockets.

The presentation is targeted at graduate students with a background in physical or computational chemistry, and it emphasizes the connections between algorithmic ideas and underlying biophysical theory.

2 Conceptual Foundations

2.1 Definition of Pockets, Cavities, and Druggability

A minimal definition of a binding site is a set of residues whose side-chains and backbone atoms contribute significantly to the binding free energy of a ligand. From a geometric perspective, computational methods often start from the protein’s solvent-accessible or solvent-excluded surface and identify:

- *Cavities*: enclosed voids within the protein.
- *Pockets*: surface indentations that can accommodate a ligand.
- *Grooves or clefts*: extended concave regions, often along interfaces.

Quantitative descriptors of a pocket include:

- Volume V and surface area A .
- Depth and exposure (fraction of buried vs. solvent-exposed surface).
- Shape descriptors (sphericity, elongation, roughness).
- Physico-chemical composition: hydrophobic/hydrophilic balance, hydrogen-bond donors/acceptors, charge distribution.

The concept of *druggability* attempts to assess whether a pocket can bind small drug-like molecules with high affinity and specificity. Druggability scores typically combine descriptors of size, shape, and physico-chemical complementarity into an empirical function calibrated against known drug-binding pockets.[6, 10] For example, many tools adopt a linear model of the form

$$D = w_V V + w_A A_{\text{hyd}} + w_P P_{\text{pol}} + w_H H_{\text{bd}} + b, \quad (1)$$

where V is pocket volume, A_{hyd} hydrophobic surface area, P_{pol} a polarity or hydrogen-bond capacity descriptor, H_{bd} counts buried hydrogen-bond donors/acceptors, and the parameters w_i, b are learned from a training set.

2.2 Thermodynamic View of Binding Sites

Binding is governed by the standard free energy of binding,

$$\Delta G_{\text{bind}} = -k_B T \ln K_a = k_B T \ln K_d, \quad (2)$$

where K_a and K_d are the association and dissociation constants, respectively, k_B is Boltzmann’s constant, and T is the absolute temperature. A favorable binding site is one for which specific protein–ligand interactions yield a large negative ΔG_{bind} , arising from:

- Enthalpic contributions: hydrogen bonds, salt bridges, dispersion interactions, metal coordination.
- Entropic contributions: desolvation of hydrophobic surfaces, restriction of conformational degrees of freedom, and release of ordered water.

At the microscopic level, a classical force-field description approximates the non-bonded interaction energy between a protein atom i and a ligand atom j as

$$E_{ij} = 4\epsilon_{ij} \left[\left(\frac{\sigma_{ij}}{r_{ij}} \right)^{12} - \left(\frac{\sigma_{ij}}{r_{ij}} \right)^6 \right] + \frac{q_i q_j}{4\pi\epsilon_0\epsilon_r r_{ij}}, \quad (3)$$

where ϵ_{ij} and σ_{ij} are Lennard–Jones parameters, r_{ij} is the interatomic distance, q_i and q_j partial charges, and ϵ_r is the effective dielectric constant. Energy-based pocket identification methods exploit such functional forms by placing probes or small fragments around the protein and computing favorable regions in terms of interaction energy or approximate binding free energy.[13, 6]

2.3 Static Versus Dynamic Views of Binding Sites

Traditionally, binding site prediction focused on static structures, typically X-ray crystal structures of the apo protein or a homolog. However, proteins are inherently dynamic, sampling ensembles of conformations according to a Boltzmann distribution,

$$p(\mathbf{R}) = \frac{1}{Z} \exp[-\beta U(\mathbf{R})], \quad \beta = \frac{1}{k_B T}, \quad (4)$$

where $\mathbf{R}$ denotes the set of atomic coordinates, $U(\mathbf{R})$ the potential energy, and Z the partition function. Cryptic and transient pockets emerge as low-population members of this ensemble: they may not be visible in any single experimental structure but can be captured by sufficiently extensive molecular dynamics (MD) simulations.[9, 4]

Two mechanistic pictures are particularly relevant:

- *Conformational selection*: the protein transiently samples an open conformation that is then selected by ligand binding.
- *Induced fit*: the ligand perturbs the protein structure upon binding, stabilizing a new conformation.

In practice, both mechanisms likely coexist, and many computational methods are implicitly tuned to detect pockets that are populated in one or both regimes.

3 Sequence-Based Binding Site Prediction

When only the amino acid sequence is available, or when structural models are uncertain, sequence-based methods provide an efficient approach to binding site prediction. These methods are particularly important in the context of large-scale proteome annotation and for targets that lack high-resolution structures.

3.1 Conservation and Motif-Based Approaches

Functionally important residues, including those forming binding sites, often exhibit higher evolutionary conservation than the surrounding surface. Given a multiple sequence alignment (MSA) of homologous sequences, per-position conservation can be quantified using the Shannon entropy,

$$H(i) = - \sum_a p_i(a) \log p_i(a), \quad (5)$$

where $p_i(a)$ is the frequency of amino acid type a at alignment position i . Positions with low entropy (high conservation) are candidates for functional sites. Mapping $H(i)$ or derived conservation scores onto a structure reveals clusters of conserved residues that often coincide with binding sites.[5]

Motif-based methods additionally search for known sequence motifs (e.g., catalytic triads, metal-binding motifs) or short linear motifs implicated in protein–protein or protein–nucleic acid interactions.[5] While these approaches can suggest regions of functional importance, they lack geometric specificity: they do not directly predict pocket shape or druggability.

3.2 Sequence-Based Machine Learning Methods

Modern sequence-based predictors treat binding site identification as a residue-level classification problem. Each residue i is represented by a feature vector $\mathbf{x}_i$ capturing:

- Local sequence context (e.g., k -mer windows).
- Evolutionary profiles from position-specific scoring matrices (PSSMs).
- Predicted structural features: secondary structure, solvent accessibility, disorder propensities.

A classifier $f(\mathbf{x}_i)$ then outputs the probability that residue i belongs to a binding site.

For instance, a logistic regression model uses

$$P(y_i = 1 \mid \mathbf{x}_i) = \sigma(\mathbf{w}^\top \mathbf{x}_i + b) = \frac{1}{1 + \exp(-\mathbf{w}^\top \mathbf{x}_i - b)}, \quad (6)$$

where $y_i \in \{0, 1\}$ indicates non-binding or binding residue, $\mathbf{w}$ and b are model parameters, and σ is the logistic function. More powerful models employ random forests, gradient-boosted trees, or deep neural networks trained to minimize cross-entropy loss over labeled residues.

Sequence-based methods are attractive due to their speed and ability to generalize across distant homologs. However, they cannot directly capture three-dimensional pocket geometry or subtle differences between alternative conformational states, limiting their resolution relative to structure-based methods.[5, 18]

4 Static Structure-Based Methods

Structure-based binding site identification assumes that a (possibly predicted) three-dimensional structure of the target or a close homolog is available. These methods may be classified into:

- Geometry-based approaches.
- Energy-based (probe/fragments) approaches.
- Template- or motif-based structure comparison.
- Hybrid and consensus methods.

4.1 Geometry-Based Pocket Detection

4.1.1 Surface and Grid Representations

The most common starting point is the solvent-excluded surface (SES) or Connolly surface, obtained by rolling a probe sphere (typically radius 1.4 Å) over the van der Waals surface of the protein. Cavities and pockets correspond to concave regions of this surface. Alternatively, methods discretize space into a three-dimensional grid and classify grid points according to whether they lie inside the protein, in solvent, or in putative cavities.[10, 18]

A simple grid-based method represents the protein by a binary occupancy function $O(\mathbf{r})$ defined on grid points $\mathbf{r}_k$,

$$O(\mathbf{r}_k) = \begin{cases} 1, & \text{if } \min_j \|\mathbf{r}_k - \mathbf{R}_j\| < r_{\text{vdW},j}, \\ 0, & \text{otherwise,} \end{cases} \quad (7)$$

where $\mathbf{R}_j$ and $r_{\text{vdW},j}$ are the coordinates and van der Waals radius of atom j . Pocket points are identified as solvent-exposed grid points adjacent to protein but enclosed from the bulk solvent; flood-fill or morphological operations are then applied to cluster these grid points into pockets.[3]

4.1.2 Alpha Shapes, Voronoi Diagrams, and Fpocket

More sophisticated geometric approaches employ alpha shapes and Voronoi diagrams. In brief, Voronoi tessellation partitions space into cells around each atom; the dual graph is the Delaunay triangulation. An *alpha complex* is then constructed by retaining only simplices (vertices, edges, tetrahedra) with circumspheres of radius less than a chosen parameter α . Cavities correspond to voids in this alpha complex.[10]

Fpocket is a widely used geometry-based method that detects pockets by clustering so-called *alpha spheres*—spheres that are tangent to four protein atoms and contain no other atomic centers.[10] Briefly, Fpocket:

1. Computes an initial set of alpha spheres from a Voronoi tessellation.
2. Filters alpha spheres based on radius (to select those likely to represent pocket-like voids).
3. Clusters nearby alpha spheres into pockets using distance-based clustering.
4. Calculates pocket descriptors (volume, polarity, hydrophobicity) and assigns a druggability score using a linear model reminiscent of Equation 1.

4.1.3 Scale-Space and Difference-of-Gaussians (DoGSite)

DoGSite constructs a 3D occupancy grid for the protein and then applies Gaussian smoothing over multiple length scales.[16] Regions that persist as density minima across scales are designated as pockets. Formally, if $\rho(\mathbf{r})$ is a smoothed occupancy field, a difference-of-Gaussians operator,

$$L(\mathbf{r}; \sigma_1, \sigma_2) = G(\mathbf{r}; \sigma_1) * \rho(\mathbf{r}) - G(\mathbf{r}; \sigma_2) * \rho(\mathbf{r}), \quad (8)$$

where G is a Gaussian kernel of width σ and $*$ denotes convolution, highlights features on a characteristic length scale. Persistent negative regions of L correspond to concavities. This multi-scale view makes DoGSite robust to noise and able to detect both small and large pockets.[16, 18]

4.2 Energy-Based and Probe-Mapping Methods

Geometry alone does not guarantee ligandability; some pockets may be geometrically prominent yet energetically unfavorable for ligand binding (e.g., highly polar, fully solvent-exposed grooves). Energy-based approaches address this by mapping interaction energies of small probe molecules onto the protein surface.

4.2.1 FTSite and FTMap

FTSite is an energy-based binding site identification method built on the FTMap fragment mapping algorithm.[13] The basic steps are:

1. A set of diverse small organic probe molecules (representing different chemotypes) is docked or sampled exhaustively over the protein surface, using a rigid-body search and fast Fourier transform (FFT)-based scoring.
2. For each probe type, low-energy poses are clustered to identify *consensus sites*.
3. Consensus clusters across different probe types are combined to yield binding hot spots; clusters of hot spots define putative ligand binding sites.

In essence, FTSite approximates the probability density of finding favorable protein–small-molecule interactions as

$$p(\mathbf{r}) \propto \sum_m \sum_{k \in \mathcal{C}_m} \exp[-\beta E_{mk}] \delta(\mathbf{r} - \mathbf{r}_{mk}), \quad (9)$$

where m indexes probe types, k configurations within cluster $\mathcal{C}_m$, and E_{mk} is the interaction energy as in Equation 3. Regions with high $p(\mathbf{r})$ across diverse probe types signal ligandable sites.[13]

4.2.2 Grid-Based Interaction Fields

An alternative is to compute grid-based interaction fields (IFPs or ΔG grids) with one or more probe types (e.g., hydrophobic carbon, hydrogen-bond donor/acceptor, cation, anion). At each grid point $\mathbf{r}_k$, one evaluates an approximate interaction free energy $\Delta G_t(\mathbf{r}_k)$ for probe type t . Sites with multiple grid points where ΔG_t is strongly favorable across complementary probe types are identified as binding hot spots. Such approaches underlie commercial tools like SiteMap.[6]

4.3 Template- and Motif-Based Structural Methods

Template-based methods exploit the observation that binding sites tend to be more conserved in structure than overall folds. Given a query protein, these algorithms:

1. Maintain a library of known binding sites described as local structural motifs (e.g., sets of residues with specific geometric and chemical patterns).
2. Superpose these motifs onto the query structure using local or global alignment of backbone or side-chain atoms.

3. Score matches by structural RMSD and chemical similarity, transferring binding site annotations from templates to the query.[5, 14]

FINDSITE, for example, uses threading to identify distant homologs in the PDB and then transfers binding site information based on conserved ligand-binding residues across structurally aligned templates.[3] SiteMotif uses graph-based pattern recognition to derive recurrent structural motifs of ligand binding sites and then searches for their occurrences in target proteins.[14]

Template-based approaches can detect subtle pockets that would be difficult to identify by geometry alone, particularly for cofactors or metal-binding sites. Their limitations include dependence on the coverage and diversity of template libraries, and potential bias toward well-studied folds.[5]

5 Machine Learning and Deep Learning for Binding Site Prediction

5.1 Traditional Machine Learning on Geometric Cavities

A natural progression from purely geometric methods is to combine geometric cavity detection with supervised learning that distinguishes true ligand-binding pockets from other surface concavities. Methods such as P2Rank and related tools follow this paradigm.[8, 18]

A typical workflow is:

1. Detect all candidate cavities using a geometric method (e.g., Fpocket, DoGSite, custom grid-based algorithms).
2. Compute cavity-level descriptors: volume, depth, hydrophobicity, electrostatics, residue composition, backbone rigidity, secondary structure, conservation scores, etc.
3. Train a classifier (e.g., random forest) on labeled examples of binding vs. non-binding cavities to predict a probability of being a ligand-binding site.

For a random forest classifier with T trees, the predicted probability that cavity c is binding is

$$P(y_c = 1 \mid \mathbf{x}_c) = \frac{1}{T} \sum_{t=1}^T h_t(\mathbf{x}_c), \quad (10)$$

where $\mathbf{x}_c$ is the feature vector and $h_t \in \{0, 1\}$ is the prediction of tree t . Decision trees implicitly partition feature space into regions of similar pockets, capturing non-linear relationships without requiring explicit feature engineering beyond cavity descriptors.

P2Rank, for instance, uses a point-based representation of the protein surface and learns to assign a ligandability score to surface points, which are then clustered into pockets.[8] Importantly, such ML-based scoring functions significantly improve the ranking of true binding sites compared to purely physics- or rule-based scoring.[18, 5]

5.2 Deep Learning on 3D Grids and Surfaces

Deep learning methods treat binding site prediction as a segmentation or detection problem on three-dimensional molecular data. Representative architectures include 3D convolutional neural networks (3D-CNNs), 3D U-Nets, and point cloud networks.[12]

5.2.1 3D Convolutional Neural Networks

In 3D-CNN approaches, the protein is embedded in a voxel grid. Each voxel contains multiple channels encoding local atomic densities or properties (e.g., atom type, partial charge, hydrophobicity). A 3D convolutional layer computes

$$y_{c'}(\mathbf{r}) = \sigma \left(\sum_c \sum_{\boldsymbol{\delta}} W_{c',c}(\boldsymbol{\delta}) x_c(\mathbf{r} + \boldsymbol{\delta}) + b_{c'} \right), \quad (11)$$

where x_c and $y_{c'}$ are input and output feature maps, $W_{c',c}$ is a learnable 3D kernel, $\boldsymbol{\delta}$ indexes positions within the kernel, and σ is a non-linear activation. Stacks of such layers capture increasingly global spatial patterns, and the network is trained to output voxel-level probabilities of being inside a binding site.[12]

Methods like DeepSite and Kalasanty exemplify this strategy, yielding high-resolution pocket segmentation and druggability estimates.[18, 12]

5.2.2 Surface and Point-Cloud Networks

Protein surfaces can be represented as point clouds or meshes with associated features (curvature, electrostatic potential, residue identity). Graph-based or point-based networks (e.g., PointNet, DGCNN variants) can then operate directly on these representations. For example, point-based architectures convert atoms into a point cloud and apply U-Net-like structures with sparse

3D convolutions to segment binding regions on protein or protein-protein interfaces.[12]

5.3 Graph Neural Networks and Geometric Deep Learning

Graph neural networks (GNNs) represent proteins as graphs with nodes corresponding to residues or atoms and edges encoding spatial proximity or chemical bonds.[12, 11] A generic message-passing layer in a GNN can be written as

$$\mathbf{h}_i^{(l+1)} = \phi \left(\mathbf{h}_i^{(l)}, \sum_{j \in \mathcal{N}(i)} \psi \left(\mathbf{h}_i^{(l)}, \mathbf{h}_j^{(l)}, \mathbf{e}_{ij} \right) \right), \quad (12)$$

where $\mathbf{h}_i^{(l)}$ is the feature vector of node i at layer l , $\mathcal{N}(i)$ is the neighborhood of i , $\mathbf{e}_{ij}$ edge features (e.g., distance, orientation), and ϕ, ψ are neural networks. After several layers, node-level outputs can be interpreted as binding propensities for residues or atoms.

GNNs are particularly appealing for cryptic and allosteric site prediction when combined with dynamic information (Section 6.4). PocketMiner, for example, is a graph neural network trained to predict residues likely to participate in cryptic pockets that open in MD simulations; it is reported to be orders of magnitude faster than MD-based screening while retaining good accuracy.[11]

5.4 Deep Learning on Sequences and Structures

Advances in protein language models (e.g., ESM, ProtBert) and structure prediction (AlphaFold2, RoseTTAFold) have enabled hybrid models that combine sequence embeddings and structural features.[12] A typical architecture might:

- Encode the sequence using a transformer-based language model to obtain residue embeddings $\mathbf{s}_i$.
- Encode the 3D structure or predicted distance maps into structural embeddings $\mathbf{t}_i$.
- Concatenate or fuse $\mathbf{s}_i$ and $\mathbf{t}_i$ through attention layers to predict residue-level binding probabilities.

The attention mechanism computes

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{d_k}}\right) \mathbf{V}, \quad (13)$$

where queries $\mathbf{Q}$, keys $\mathbf{K}$, and values $\mathbf{V}$ are linear projections of $\{\mathbf{s}_i, \mathbf{t}_i\}$, and d_k is the key dimension. Such models can naturally learn long-range couplings between residues, which are essential for recognizing distributed allosteric networks and extended binding sites.[12]

6 Dynamics-Based Methods and Cryptic/Transient Pockets

Static approaches capture only pockets present in a single conformation. To reveal cryptic and transient pockets, it is necessary to sample protein dynamics. This section discusses MD-based approaches, enhanced sampling techniques, and specialized methods for cryptic pocket discovery.

6.1 Classical Molecular Dynamics and Pocket Ensembles

Standard MD simulations generate time series of atomic coordinates $\{\mathbf{R}(t)\}$ by integrating Newton’s equations of motion under a molecular mechanics force field. Pocket analysis tools such as MDpocket, TRAPP, or in-house scripts can then detect cavities in each snapshot and accumulate statistics over the trajectory.[9, 4]

Let $\mathcal{P}_s$ denote the set of grid points assigned to pocket p in snapshot s . The occupancy of grid point k by pocket p over a trajectory of S snapshots is

$$O_p(k) = \frac{1}{S} \sum_{s=1}^S \mathbf{1}\{k \in \mathcal{P}_s\}, \quad (14)$$

where $\mathbf{1}\{\cdot\}$ is the indicator function. High values of $O_p(k)$ indicate persistent pockets, while intermediate values reflect transient pockets that open and close during the simulation.

Pocket volume distributions over time, $P_p(V)$, provide insight into the dynamics of pocket opening. For cryptic pockets, one often observes a low baseline volume punctuated by rare opening events with significantly larger volumes.[9]

6.2 Cosolvent and Mixed-Solvent MD

Mixed-solvent MD (MSMD), also called cosolvent MD, enhances sampling of binding hot spots by adding small organic probes (e.g., isopropanol, acetonitrile, benzene) at low molar fractions to the aqueous solvent.[9, 4] Probes preferentially accumulate in regions of the protein that can favorably accommodate hydrophobic or polar groups, thereby highlighting potential binding sites.

The probe density for species t at grid point $\mathbf{r}_k$ is estimated as

$$\rho_t(\mathbf{r}_k) = \frac{1}{S} \sum_{s=1}^S \sum_{m \in \mathcal{M}_t(s)} \mathbf{1}\{\|\mathbf{r}_k - \mathbf{r}_m(s)\| < r_{\text{cut}}\}, \quad (15)$$

where $\mathcal{M}_t(s)$ is the set of probe molecules of type t at snapshot s , $\mathbf{r}_m(s)$ their coordinates, and r_{cut} a cutoff distance. Hot spots are defined as regions where $\rho_t(\mathbf{r}_k)$ exceeds a threshold for one or more probe types.

MSMD provides both spatial (location of hot spots) and dynamic (occupation times, opening frequencies) information. However, cryptic pockets that require large backbone rearrangements may still be rare events even with cosolvents, motivating the use of enhanced sampling (Section 6.3) and recent developments such as CrypToth, which integrate MSMD with topological data analysis.[7, 4]

6.3 Enhanced Sampling Methods

Enhanced sampling techniques aim to accelerate exploration of conformational space beyond timescales accessible to straightforward MD. For cryptic pocket discovery, the relevant rare events are transitions between closed and open pocket conformations. Methods include accelerated MD, metadynamics, Hamiltonian scaling (e.g., SWISH, SWISH-X), and weighted-ensemble MD.[9, 2, 17]

6.3.1 Accelerated MD and Ligand-Mapping MD

Accelerated MD (aMD) modifies the potential energy surface by adding a bias potential that raises energy wells below a threshold E ,

$$U'(\mathbf{R}) = \begin{cases} U(\mathbf{R}) + \Delta U(\mathbf{R}), & U(\mathbf{R}) < E, \\ U(\mathbf{R}), & U(\mathbf{R}) \geq E, \end{cases} \quad (16)$$

with

$$\Delta U(\mathbf{R}) = \frac{(E - U(\mathbf{R}))^2}{\alpha + (E - U(\mathbf{R}))}, \quad (17)$$

where α controls the degree of acceleration. This flattens energy wells, enhancing barrier crossing and conformational transitions. Ligand-mapping MD (LMMD) combines MD with small probe ligands to identify binding hot spots dynamically. Accelerated ligand-mapping MD (aLMMD) merges aMD with LMMD to efficiently detect deeply buried cryptic pockets and occluded binding sites.[15]

6.3.2 Metadynamics and Collective Variables

Metadynamics introduces a history-dependent bias potential along a chosen set of collective variables (CVs), such as pocket volume V_p or distances between gating residues. The bias at time t is

$$V_{\text{bias}}(\mathbf{s}, t) = \sum_{\tau < t} w \exp \left[-\frac{\|\mathbf{s} - \mathbf{s}(\tau)\|^2}{2\sigma^2} \right], \quad (18)$$

where $\mathbf{s}$ are the CVs, w is the Gaussian height, and σ their width. As bias accumulates in previously sampled regions, the system is driven to explore new configurations, potentially opening cryptic pockets.[9] Choice of CVs is crucial; for cryptic pockets, useful CVs may encode inter-residue distances defining a gate, pocket volume, or coordinates of cavity-filling side-chains.[4]

6.3.3 Hamiltonian Scaling and SWISH / SWISH-X

The SWISH (Sampling Water Interfaces through Scaled Hamiltonians) protocol scales protein–water interactions to facilitate opening of hydrophobic pockets by effectively making water less favorable at certain interfaces.[9] In SWISH-X, SWISH is combined with On-the-fly Probability Enhanced Sampling (OPES) MultiThermal to improve sampling across a spectrum of effective temperatures.[2] Conceptually, these methods explore an ensemble of modified Hamiltonians,

$$U_{\lambda}(\mathbf{R}) = U_{\text{protein}}(\mathbf{R}) + \lambda U_{\text{protein–water}}(\mathbf{R}) + U_{\text{water}}(\mathbf{R}), \quad (19)$$

with $\lambda < 1$ temporarily weakening protein–water interactions to encourage dehydration and pocket opening; proper reweighting recovers observables of the original system.[2, 17]

SWISH-X has been shown to outperform standard MSMD and SWISH in opening cryptic pockets at protein–protein interfaces and isolated proteins, revealing cavities inaccessible to unbiased simulations in similar computational time.[2]

6.3.4 Weighted-Ensemble MD for Pocket Opening

Weighted-ensemble (WE) MD methods focus computational effort on rare transitions by partitioning the CV space (e.g., pocket volume) into bins and spawning or pruning trajectories based on their weights. This enables direct estimation of transition rates between closed and open pocket states and facilitates mapping of pocket opening pathways.[4, 17]

6.4 Cryptic Pocket Detection: Methods and Theory

6.4.1 Definition and Challenges

Cryptic pockets are defined as ligand-binding pockets that are absent or occluded in the apo structure but appear in ligand-bound structures or rare conformations.[9, 1, 4] From a thermodynamic perspective, the open pocket state is often higher in free energy than the closed state in the absence of ligand,

$$\Delta G_{\text{open-closed}} = -k_{\text{B}}T \ln \frac{p_{\text{open}}}{p_{\text{closed}}} > 0, \quad (20)$$

leading to low equilibrium occupancy p_{open} . Ligand binding stabilizes the open state, shifting the equilibrium via conformational selection. The rarity and structural heterogeneity of cryptic pockets pose challenges for both MD sampling and ML model training.[9, 4]

6.4.2 MD-Based Approaches and Markov State Models

Extended MD simulations analyzed with Markov state models (MSMs) have been used to dissect cryptic pocket opening mechanisms and kinetics.[9, 4] In an MSM, the conformational space is discretized into states (e.g., microstates based on structural clustering), and a transition probability matrix $\mathbf{T}(\tau)$ is estimated for lag time τ ,

$$T_{ij}(\tau) = P(\text{state } j \text{ at } t + \tau \mid \text{state } i \text{ at } t). \quad (21)$$

Diagonalization of $\mathbf{T}$ yields implied timescales for slow processes, including pocket opening and closing. MSMs can reveal which conformational pathways lead to cryptic pocket formation and identify gating residues whose motions control access.[9]

Enhanced-sampling simulations such as SWISH, SWISH-X, and aLMMD can be analyzed similarly, though careful reweighting is required to extract unbiased kinetics.[15, 2, 17]

6.4.3 Data-Centric and Exposon-Based Methods

Data-centric approaches, such as the “Will the Real Cryptic Pocket Please Stand Out?” framework, leverage large MD datasets and unsupervised learning to identify correlated solvent exposure changes (“exposons”) that signal cryptic site formation.[1] Here, per-residue solvent-accessible surface area (SASA) time series $A_i(t)$ are clustered into groups of residues with correlated exposure dynamics; clusters that transiently increase SASA together often delineate cryptic pockets.

Formally, one may compute a correlation matrix C_{ij} of SASA fluctuations,

$$C_{ij} = \frac{\langle \delta A_i(t) \delta A_j(t) \rangle}{\sqrt{\langle \delta A_i(t)^2 \rangle \langle \delta A_j(t)^2 \rangle}}, \quad (22)$$

cluster residues based on C_{ij} , and then analyze structural mapping of these clusters to identify putative cryptic pockets.[1, 9]

6.4.4 Machine-Learning Methods for Cryptic Pockets

Several ML-based methods specialize in cryptic pocket detection.[4, 17]

- **CryptoSite**: an SVM-based classifier that distinguishes cryptic-site residues based on differences between apo and holo structures and local environment descriptors.
- **TACTICS**: a random forest trained on reconstructed CryptoSite data and ConCavity features, able to rank cryptic pockets and assess drug-gability, including via fragment docking.
- **PocketMiner**: a graph neural network that predicts residues likely to form cryptic pockets using features from static structures and labels derived from MD trajectories.[11]
- **Other deep models**: networks (e.g. Ssnet) that use backbone curvature and torsion descriptors or 3D grids, interpreted via attribution methods (Grad-CAM, integrated gradients), to highlight cryptic sites based on structural features alone.[4]

These models abstract over specific MD details and can be applied proteome-wide to prioritize proteins and regions for more expensive simulations or experimental study.[11, 4]

6.4.5 CrypToth: MSMD with Topological Data Analysis

CrypToth is a recent method that couples MSMD with topological data analysis (TDA), specifically persistent homology, to distinguish cryptic sites from generic hot spots.[7] After running MSMD with multiple probe types, probe occupancy maps are converted into scalar fields on the protein surface. Persistent homology tracks the birth and death of topological features (connected components, holes) as an occupancy threshold is varied. Hot spots whose topological signatures persist over a broad range of thresholds and correlate with conformational flexibility are ranked highly as cryptic pocket candidates.[7, 4]

Mathematically, persistent homology constructs a filtration of simplicial complexes $\{K_\alpha\}$ over a filtration parameter α (e.g., occupancy threshold) and computes Betti numbers $\beta_k(\alpha)$ that count k -dimensional holes. The lifetime of features in a persistence diagram provides a measure of their robustness. CrypToth integrates these lifetimes with energetic and dynamic metrics to score sites.[7]

7 Validation of Predicted Binding Sites

Prediction alone is insufficient; computationally identified sites must be validated and prioritized for experimental work. Validation strategies span internal consistency checks, benchmarking against known datasets, energetic analyses, evolutionary evidence, and direct experimental measurements.

7.1 Benchmarking and Performance Metrics

Comparative evaluations of binding site predictors have standardized performance metrics and benchmark datasets.[5, 18] Given a protein with one or more known binding sites and a predictor that outputs a ranked list of candidate pockets, common metrics include:

- **Top- N recall:** fraction of proteins for which at least one true binding site overlaps with the top N predicted sites.
- **Precision and recall** at residue or grid-point level:

$$\text{Precision} = \frac{TP}{TP + FP}, \quad \text{Recall} = \frac{TP}{TP + FN}, \quad (23)$$

where TP are true positives (correctly predicted binding residues), FP false positives, and FN false negatives.

- **F1-score:**

$$F_1 = 2 \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}. \quad (24)$$

When introducing new methods, it is essential to benchmark against curated datasets (e.g., scPDB, PDBbind-based sets) and report such metrics. Reviews by Zhang *et al.* and others provide overviews of widely used datasets and evaluation protocols.[18]

7.2 Cross-Method Consistency and Consensus

A pragmatic validation step is to compare results across multiple methods:

- Agreement between geometry-based tools (Fpocket, DoGSite), energy-based mapping (FTSite, MSMD), and ML predictors (P2Rank, DL-based) supports the significance of a site.[10, 16, 8]
- High cryptic-pocket scores from PocketMiner or CryptoSite that coincide with pockets seen in enhanced MD or MSMD simulations bolster confidence.[11, 9]

Consensus strategies can be formalized via ensemble averaging or voting schemes. For example, given method-specific scores $s_m(p)$ for pocket p , a consensus score might be

$$S_{\text{cons}}(p) = \sum_m \alpha_m \hat{s}_m(p), \quad (25)$$

where $\hat{s}_m$ are normalized scores and α_m weights chosen based on prior performance or cross-validation.

7.3 Druggability and Energetic Validation

Druggability assessment combines geometric and energetic information to estimate whether a pocket is likely to bind small molecules with high affinity.[6, 10] Beyond empirical scoring (Equation 1), one can perform:

- **Docking of fragment libraries** into the predicted site and analyze enrichment of high-scoring poses.
- **Binding free energy estimates** for representative ligands using endpoint methods such as MM/GBSA:

$$\Delta G_{\text{bind}}^{\text{MM/GBSA}} = \langle E_{\text{complex}} \rangle - \langle E_{\text{protein}} \rangle - \langle E_{\text{ligand}} \rangle + \Delta G_{\text{solv}} - T\Delta S, \quad (26)$$

where $\langle E \rangle$ are average gas-phase energies and ΔG_{solv} solvation contributions (often from a generalized Born model).

- **Alchemical free energy calculations** for ligand binding to different pockets or different conformational states (e.g., open vs. closed cryptic pocket) to quantify relative affinities.[9, 17]

For cryptic or transient pockets, a key validation is whether ligands designed or docked into the open pocket remain bound and stable in unbiased MD simulations initiated from those complexes, indicating a realistic environment and feasible access path.[1, 9]

7.4 Evolutionary and Coevolutionary Signals

Mapping evolutionary conservation and coevolutionary couplings onto predicted pockets provides orthogonal evidence of functional relevance.[5] Highly conserved residues located within a predicted pocket suggest functional or regulatory roles. Coevolutionary analysis (e.g., direct coupling analysis) can identify networks of residues that co-vary, many of which participate in allosteric pathways connecting distant sites to functional regions.

For cryptic pockets, conservation may be weaker, as these sites often confer specificity or allosteric regulation unique to subfamilies. Nonetheless, clusters of moderately conserved residues near cryptic pockets can signal evolutionary tuning of allosteric control.[4, 17]

7.5 Experimental Validation of Binding Sites

Ultimately, the strongest validation comes from experiments. Common strategies include:

- **X-ray crystallography or cryo-EM** of the protein in complex with ligands designed or identified to target the predicted site. For cryptic pockets, this can require co-crystallization with stabilizing ligands or soaking under conditions that favor pocket opening.[9]
- **NMR chemical shift perturbation (CSP)** and **paramagnetic relaxation enhancement (PRE)** to map ligand-binding epitopes and detect conformational shifts upon binding.
- **Hydrogen–deuterium exchange mass spectrometry (HDX-MS)** to measure local changes in backbone amide protection upon ligand binding, highlighting regions that become buried or structurally stabilized.[9, 4]

- **Site-directed mutagenesis** of key pocket residues (e.g., gating side-chains) to assess effects on ligand binding affinity and functional read-outs (enzyme activity, signaling).

For cryptic pockets, experimental validation often combines conformational probes with functional assays. For example, cryptic pockets discovered by MD and targeted by small molecules have yielded both positive and negative allosteric modulators, demonstrating that binding at cryptic sites can tune function in diverse ways.[1, 4]

7.6 Specific Considerations for Cryptic and Transient Pockets

Validation of cryptic and transient pockets involves additional challenges:[9, 17]

- **Reproducibility across simulations:** independent MD or enhanced-sampling runs should repeatedly sample the same pocket with comparable volumes and structural features.
- **Opening frequency and lifetime:** MSM or WE analyses can quantify the equilibrium probability and mean first passage times for pocket opening, providing a measure of how accessible the site is in physiological conditions.
- **Ligand-induced stabilization:** ligands designed for the cryptic pocket should stabilize the open conformation and shift the equilibrium, observable via HDX-MS, NMR, or thermal shift assays.
- **Allosteric coupling:** perturbations at the cryptic pocket (mutations, ligand binding) should produce measurable changes in activity or interactions at the orthosteric site or other functional regions.[4]

Recent reviews synthesize the state of the art in cryptic pocket detection and validation, emphasizing that integrating physics-based simulations (MSMD, SWISH-X, aLMMD, WE) with AI-based pre-screening (Pocket-Miner, CryptoSite, TACTICS) and targeted experiments provides the most robust pipeline.[9, 4, 17]

8 Practical Workflows for Binding Site Identification

8.1 Static Structure-Based Workflow

For a newly determined structure (experimental or high-confidence predicted):

1. **Preprocessing:** assign protonation states, add missing side-chains or loops, and relax the structure with restrained MD or minimization.
2. **Geometry-based screening:** run Fpocket, DoGSite, and related tools to identify surface pockets and obtain initial druggability scores.[10, 16]
3. **Energy-based mapping:** apply FTMap/FTSite or equivalent probe-mapping to confirm energetically favorable hot spots within geometric pockets.[13]
4. **ML-based scoring:** use P2Rank or similar ML predictors to re-rank pockets by ligandability.[8]
5. **Template search:** compare against binding-site templates (e.g., Site-Motif, FINDSITE) to detect conserved motifs or cofactor sites.[3, 14]
6. **Consensus and prioritization:** integrate scores into a consensus ranking; select top candidates for docking or experimental follow-up.

8.2 Dynamics-Enhanced Workflow for Cryptic/Transient Pockets

When cryptic or transient pockets are suspected (e.g., in protein–protein interfaces, “undruggable” targets, or kinases with known allosteric regulation):

1. **Fast pre-screening:** run PocketMiner or related ML tools to identify residues with high cryptic-pocket propensity.[11, 4]
2. **Classical MD:** perform multiple independent MD simulations (or one long trajectory) and analyze pocket dynamics using MDpocket or similar, computing occupancy and volume distributions (Equation 14).[9]
3. **MSMD:** run mixed-solvent MD with diverse probes to map hot spots and correlate with dynamically opening pockets.[9, 7]

4. **Enhanced sampling:** for promising but rarely sampled pockets, apply SWISH-X, aLMMD, or WE MD to enhance sampling of opening events and estimate kinetics.[15, 2, 17]
5. **Post hoc ML/TDA analysis:** apply methods like CrypToth to rank MSMD hot spots and quantify cryptic pocket robustness using persistent homology and dynamic flexibility metrics.[7]
6. **Ligand design and validation:** design fragments or small molecules targeting the cryptic pocket, validate via docking, MD stability, and experimental assays (CSP, HDX-MS, crystallography).[9, 4]

9 Summary and Outlook

Computational identification of binding sites on proteins has matured from heuristic geometry-based algorithms into a rich ecosystem of physically grounded, statistically principled, and machine-learning-driven methods. Geometry-based and energy-based approaches provide fast and interpretable tools for detecting classical surface pockets and ranking their druggability. Template-based strategies leverage the evolutionary reuse of structural motifs to recognize binding sites even in proteins with low global homology. Machine-learning and deep-learning methods harness large datasets of known complexes to learn complex, non-linear patterns of ligandability in sequence and three-dimensional space.[5, 18, 12]

In parallel, advances in MD simulations and enhanced sampling, combined with data-centric and AI-guided analyses, have transformed the understanding and exploitation of cryptic and transient pockets. Such sites greatly expand the druggable proteome, particularly for historically “undruggable” targets lacking obvious orthosteric pockets. Methods like MSMD, SWISH-X, aLMMD, weighted-ensemble MD, PocketMiner, and CrypToth exemplify the trend toward tightly integrating physics-based dynamics with sophisticated analysis and ML models.[9, 11, 7, 17]

Future directions include:

- Tighter integration of binding site prediction with generative models for ligand design, allowing end-to-end pipelines from target sequence to candidate ligands.
- Systematic incorporation of experimental and biophysical data (HDX-MS, NMR relaxation, cryo-EM ensembles) as constraints or training signals for pocket detection.

- Uncertainty quantification in predictions, enabling principled ranking of pockets for experimental follow-up.
- Application of quantum chemical methods and quantum computing for high-accuracy energetic characterization of challenging pockets, including covalent and metalloprotein sites.

For graduate students entering the field, understanding the theoretical underpinnings of these methods—from geometric representations and statistical mechanics to ML architectures and sampling theory—is essential for both effective application and further methodological innovation.

References

- [1] Rommie E Amaro. Will the real cryptic pocket please stand out? *Bio-physical Journal*, 116(5):753–754, 2019.
- [2] Alberto Borsatto, Eleonora Gianquinto, Valerio Rizzi, and Francesco Luigi Gervasio. Swish-x, an expanded approach to detect cryptic pockets in proteins and at protein–protein interfaces. *Journal of Chemical Theory and Computation*, 20(8):3335–3348, 2024.
- [3] Michal Brylinski and Jeffrey Skolnick. A threading-based method (find-site) for ligand-binding site prediction and functional annotation. *Proceedings of the National Academy of sciences*, 105(1):129–134, 2008.
- [4] Dorota Gašparíková, Rupesh Chikhale, Jason Cole, and Ehmke Pohl. Recent computational advances in the identification of cryptic binding sites for drug discovery. *Bioinformatics Advances*, 5(1):vbaf156, 2025.
- [5] Dario Ghersi and Roberto Sanchez. Beyond structural genomics: computational approaches for the identification of ligand binding sites in protein structures. *Journal of structural and functional genomics*, 12(2):109–117, 2011.
- [6] Thomas A Halgren. Identifying and characterizing binding sites and assessing druggability. *Journal of chemical information and modeling*, 49(2):377–389, 2009.
- [7] Jun Koseki, Chie Motonou, Keisuke Yanagisawa, Genki Kudo, Ryunosuke Yoshino, Takatsugu Hirokawa, and Kenichiro Imai. Cryptoth: Cryptic pocket detection through mixed-solvent molecular dynamics simulations-based topological data analysis. *Journal of Chemical Information and Modeling*, 65(11):5567–5575, 2025.

- [8] Radoslav Krivák and David Hoksza. P2rank: machine learning based tool for rapid and accurate prediction of ligand binding sites from protein structure. *Journal of cheminformatics*, 10(1):39, 2018.
- [9] Antonija Kuzmanic, Gregory R Bowman, Jordi Juarez-Jimenez, Julien Michel, and Francesco L Gervasio. Investigating cryptic binding sites by molecular dynamics simulations. *Accounts of chemical research*, 53(3):654–661, 2020.
- [10] Vincent Le Guilloux, Peter Schmidtke, and Pierre Tuffery. Fpocket: an open source platform for ligand pocket detection. *BMC bioinformatics*, 10(1):168, 2009.
- [11] Artur Meller, Michael D Ward, Jonathan H Borowsky, Jeffrey M Lotthammer, Meghana Kshirsagar, Felipe Oviedo, Juan Lavista Ferres, and Gregory Bowman. Predicting the locations of cryptic pockets from single protein structures using the pocketminer graph neural network. *Biophysical journal*, 122(3):445a, 2023.
- [12] Minjie Mou, Zhichao Zhang, Ziqi Pan, and Feng Zhu. Deep learning for predicting biomolecular binding sites of proteins. *Research*, 8:0615, 2025.
- [13] Chi-Ho Ngan, David R Hall, Brandon Zerbe, Laurie E Grove, Dima Kozakov, and Sandor Vajda. Ftsite: high accuracy detection of ligand binding sites on unbound protein structures. *Bioinformatics*, 28(2):286–287, 2012.
- [14] Santhosh Sankar and Nagasuma Chandra. Sitemotif: A graph-based algorithm for deriving structural motifs in protein ligand binding sites. *PLoS Computational Biology*, 18(2):e1009901, 2022.
- [15] Justin Tze-Yang Ng and Yaw Sing Tan. Accelerated ligand-mapping molecular dynamics simulations for the detection of recalcitrant cryptic pockets and occluded binding sites. *Journal of Chemical Theory and Computation*, 18(3):1969–1981, 2022.
- [16] Andrea Volkamer, Axel Griewel, Thomas Grombacher, and Matthias Rarey. Analyzing the topology of active sites: on the prediction of pockets and subpockets. *Journal of chemical information and modeling*, 50(11):2041–2052, 2010.

- [17] Si Zhang and Gregory R Bowman. Decrypting cryptic pockets with physics-based simulations and artificial intelligence. *Current Opinion in Structural Biology*, 96:103215, 2026.
- [18] Jingtian Zhao, Yang Cao, and Le Zhang. Exploring the computational methods for protein-ligand binding site prediction. *Computational and structural biotechnology journal*, 18:417–426, 2020.